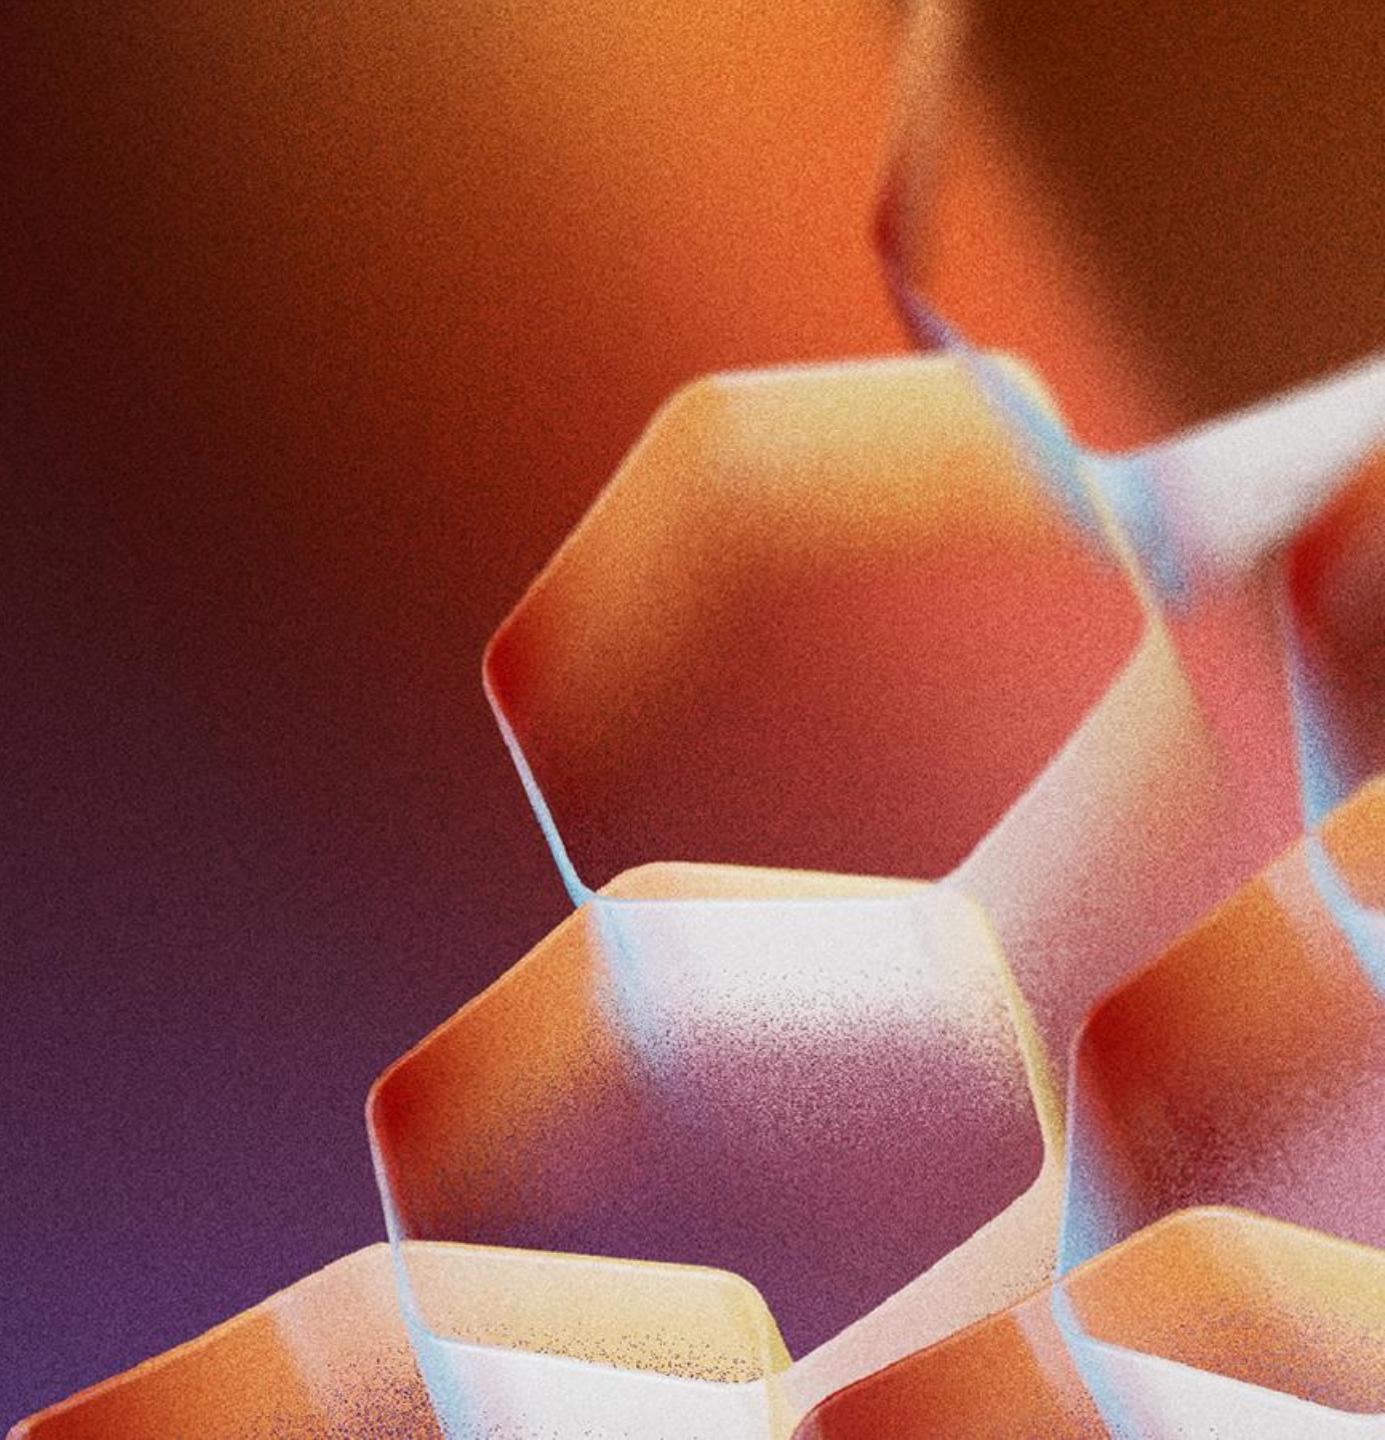




# Microsoft Build





# KBYG: AI Expert Topic Readiness call

May 13, 2025

Laura Wolf, Expert Recruit Lead (ERL)



# Agenda

- Expert Meetup overview
- Logistics
- Product announcements – Jenn Cockrell
- AI Expert Leads – Laura Wolf
- Resources
  - GitHub Repo – Nick Brady
  - Feedback – Marc Baiza
  - Gartner reviews program – Michal Toiba
- Focus area deep dives – Expert Leads
- Community – Jonathan Jones

# Microsoft Build Expert Meetup Key Players



**Dawn Conner**  
Hub Experiences  
& Expert Meetup  
Lead



**Jonathan Jones**  
Community Lead  
& Expert Meetup  
Support



**Shawndell  
Elfring**  
Hub Experiences  
& Expert Meetup  
PM



**Paige Boggie**  
Expert Staffing  
PM



**Laura Wolf**  
AI Expert Recruit  
Lead

# Expert Meetup Program Overview

- For attendee:** An opportunity for attendees engage in 1:1 or 1:few conversations and demos with Microsoft Experts/SMEs to address technical questions. This program aims to address technical questions, enhance product and solution knowledge, remove customer business roadblocks, and support company news and announcements through interactive discussions.
- For Microsoft:** Connecting with attendees gives us direct insight into our customers allowing us to track verbatim feedback, share our passion for our products and services, connect with industry leaders, and help our attendees build relationships and collaborations both inside Microsoft and with the community.

## Expert Meetup Details

- Expert Meetups will be located in Hub
- Day 1-3 (May 19-21)
- This is a persistent 'always on'
- Dedicated spaces to facilitate 1:1 and 1:few conversations
- Organized by topic and sub-topics (focus areas)
- The emphasis is on facilitating conversations with attendees

Topic	Expert Recruit Lead (ERL)
AI, Copilot & Agents	Laura Wolf & Dona Sarkar
Developer Tools & .NET	Bruno Capuano
Cloud Platform	Vinicius Apolinario
Data & Analytics	Paul DeCarlo
Security	Mykaela Gardner
Windows	Arnaud Montembault
Copilot+ PC	Jacob Rhoades
GitHub	Anisha Pindoria
Code for Good	Amy Hua

## Expert Roles & Responsibilities

### Pre-Event

- Join our Expert Know Before You Go call, hosted on May 7
- Join ERL led Topic Readiness Call, hosted May 5-9
- Accept shift calendar holds, sent April 29
- If your scheduling availability changes, work with your ERL to update shift assignments as needed

### Onsite Shift Responsibilities

- Check-in 30 mins before the start of each shift at the designated Expert Shift Check-in Desk
- Keep personal items to a minimum and place them in the storage cabinet during your shift
- Remain at your assigned station during the entirety of your shift
- Do not leave your station unattended; coordinate break coverage with your ERL or other experts

### Attendee Engagement

- Engage with attendees and answer questions related to your topic of expertise
- Be familiar with sessions and partners related to your topic of expertise and make recommendations

### Content

- If you plan to demo, you may use any content that your ERL has pre-loaded on the station device
- You may also bring materials on a USB stick or use your own device

### Onsite Escalations

- Reach out to your ERL with any questions or onsite escalations as needed

# The Hub

## SCC Arch Level 4

### Microsoft Experiences

- Expert Meetup (Topic Areas)
- GitHub
- Microsoft Learn
- Pegasus Startups
- Code for Good (AI for Good)
- 50<sup>th</sup> Anniversary Gallery

### Community

- Networking Tables

### Featured Partners

### Demo Theaters

### Breakouts

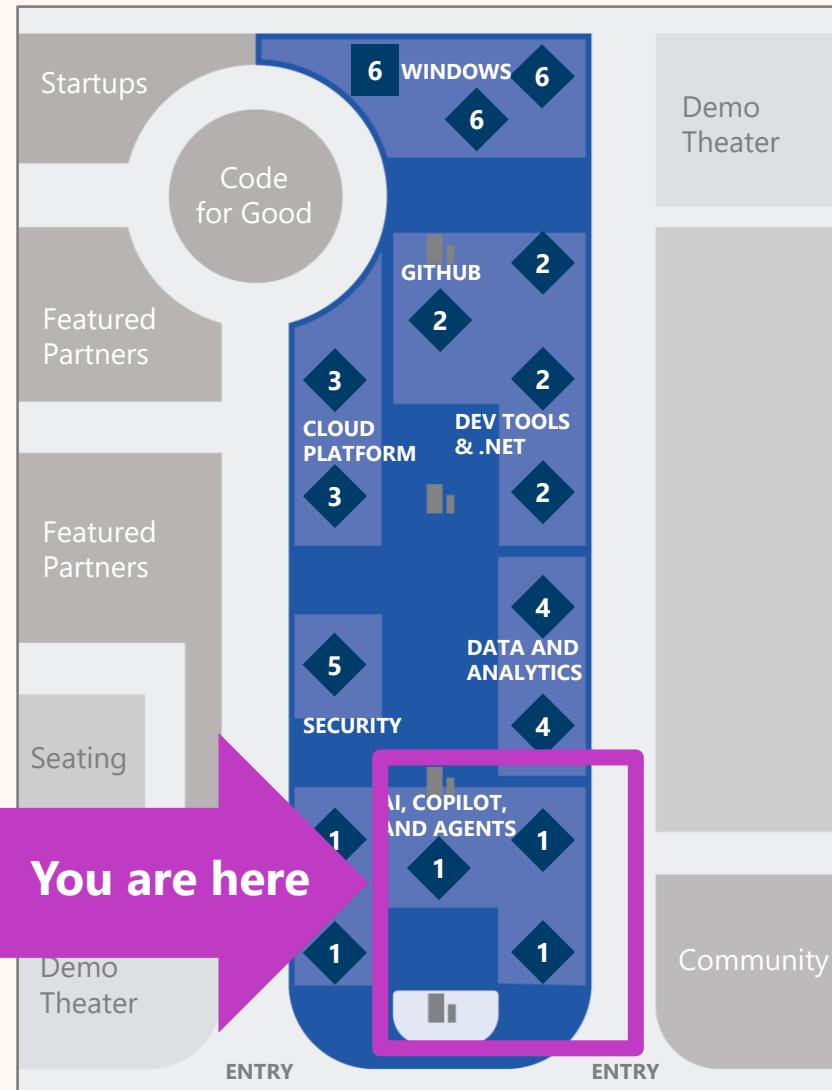
### Event Essentials

- Food & Beverage
- Badge pickup



# Expert Meetup Focus Areas

Topics	Focus Areas	Recruit Lead	CSA Lead
1 AI, Copilot, and Agents	<ul style="list-style-type: none"> <li>Azure AI Foundry</li> <li>Azure AI Foundry Models</li> <li>Azure OpenAI in Foundry Models</li> <li>Azure AI Foundry Agent Service</li> <li>Azure AI Search Agentic Retrieval</li> <li>Model Customization</li> </ul>	<ul style="list-style-type: none"> <li>Foundry Local</li> <li>Azure AI Document Intelligence</li> <li>Azure AI Tools</li> <li>Azure AI Content Safety</li> <li>Azure AI Foundry Observability</li> <li>AI Skilling</li> </ul>	Laura Wolf Kavita Makdani
	<ul style="list-style-type: none"> <li>Microsoft Copilot Studio</li> <li>Microsoft 365 Copilot</li> <li>Copilot Control Systems</li> <li>Agents and MCP servers for Dynamics 365</li> </ul>	<ul style="list-style-type: none"> <li>Agent in Microsoft 365</li> <li>Microsoft Teams</li> <li>Microsoft SharePoint</li> <li>Agent Store</li> <li>Microsoft Security Copilot agents</li> </ul>	
2 Developer Tools and .NET	<ul style="list-style-type: none"> <li>.NET and C#</li> <li>.NET Aspire for Cloud Apps</li> <li>Visual Studio</li> <li>.NET AI Development</li> <li>Visual Studio Code</li> <li>Well-architected SaaS</li> </ul>	<ul style="list-style-type: none"> <li>Agentic Developer Tools</li> <li>Agentic DevOps</li> <li>Model Context Protocol (MCP)</li> <li>Microsoft Dev Box and DevTest Labs</li> <li>Power Platform</li> <li>Agents with Dataverse</li> </ul>	Bruno Capuano Jeff Medford, Kavita Makdani
	<ul style="list-style-type: none"> <li>GitHub Copilot</li> <li>AI-Native</li> </ul>	<ul style="list-style-type: none"> <li>Enterprise Scale</li> <li>Code Security</li> </ul>	
3 Cloud Platform	<ul style="list-style-type: none"> <li>Cloud Native and Azure Kubernetes Service</li> <li>Azure App Service</li> <li>Serverless</li> <li>Azure API Management and Azure Logic Apps</li> </ul>	<ul style="list-style-type: none"> <li>Azure Arc and Azure Monitor</li> <li>Open Source Software and Linux on Azure</li> <li>Azure Infrastructure</li> <li>AI-assisted App modernization</li> </ul>	Vinicius Apolinario Kavita Makdani
4 Data and Analytics	<ul style="list-style-type: none"> <li>Microsoft Fabric</li> <li>Microsoft Power BI</li> <li>Databases in Microsoft Fabric</li> <li>Azure Cosmos DB</li> </ul>	<ul style="list-style-type: none"> <li>Azure SQL DB &amp; SQL Server</li> <li>Azure Databases for PostgreSQL &amp; MySQL</li> <li>Azure Managed Redis</li> <li>Azure Databricks</li> </ul>	Paul DeCarlo Kavita Makdani
5 Security	<ul style="list-style-type: none"> <li>AI red teaming</li> <li>Identity and network security</li> </ul>	<ul style="list-style-type: none"> <li>Data Security for AI</li> <li>Cloud security</li> </ul>	Mykaela Gardner Candice Hargett
6 Windows	<ul style="list-style-type: none"> <li>Windows Developer experiences</li> <li>Windows AI Foundry</li> <li>Arm Development on Windows</li> <li>Windows App Development</li> </ul>	<ul style="list-style-type: none"> <li>MCP &amp; App Actions on Windows</li> <li>Security and Privacy</li> <li>Windows Insider Program</li> <li>Windows 3P apps and Microsoft Store</li> </ul>	Arnaud Montembault Divya Venkataramu
	<ul style="list-style-type: none"> <li>Copilot+ PC - Hands-on learning with Surface</li> <li>Windows AI Developer Workstations</li> </ul>		
		Jacob Rhoades Arnaud Montembault	



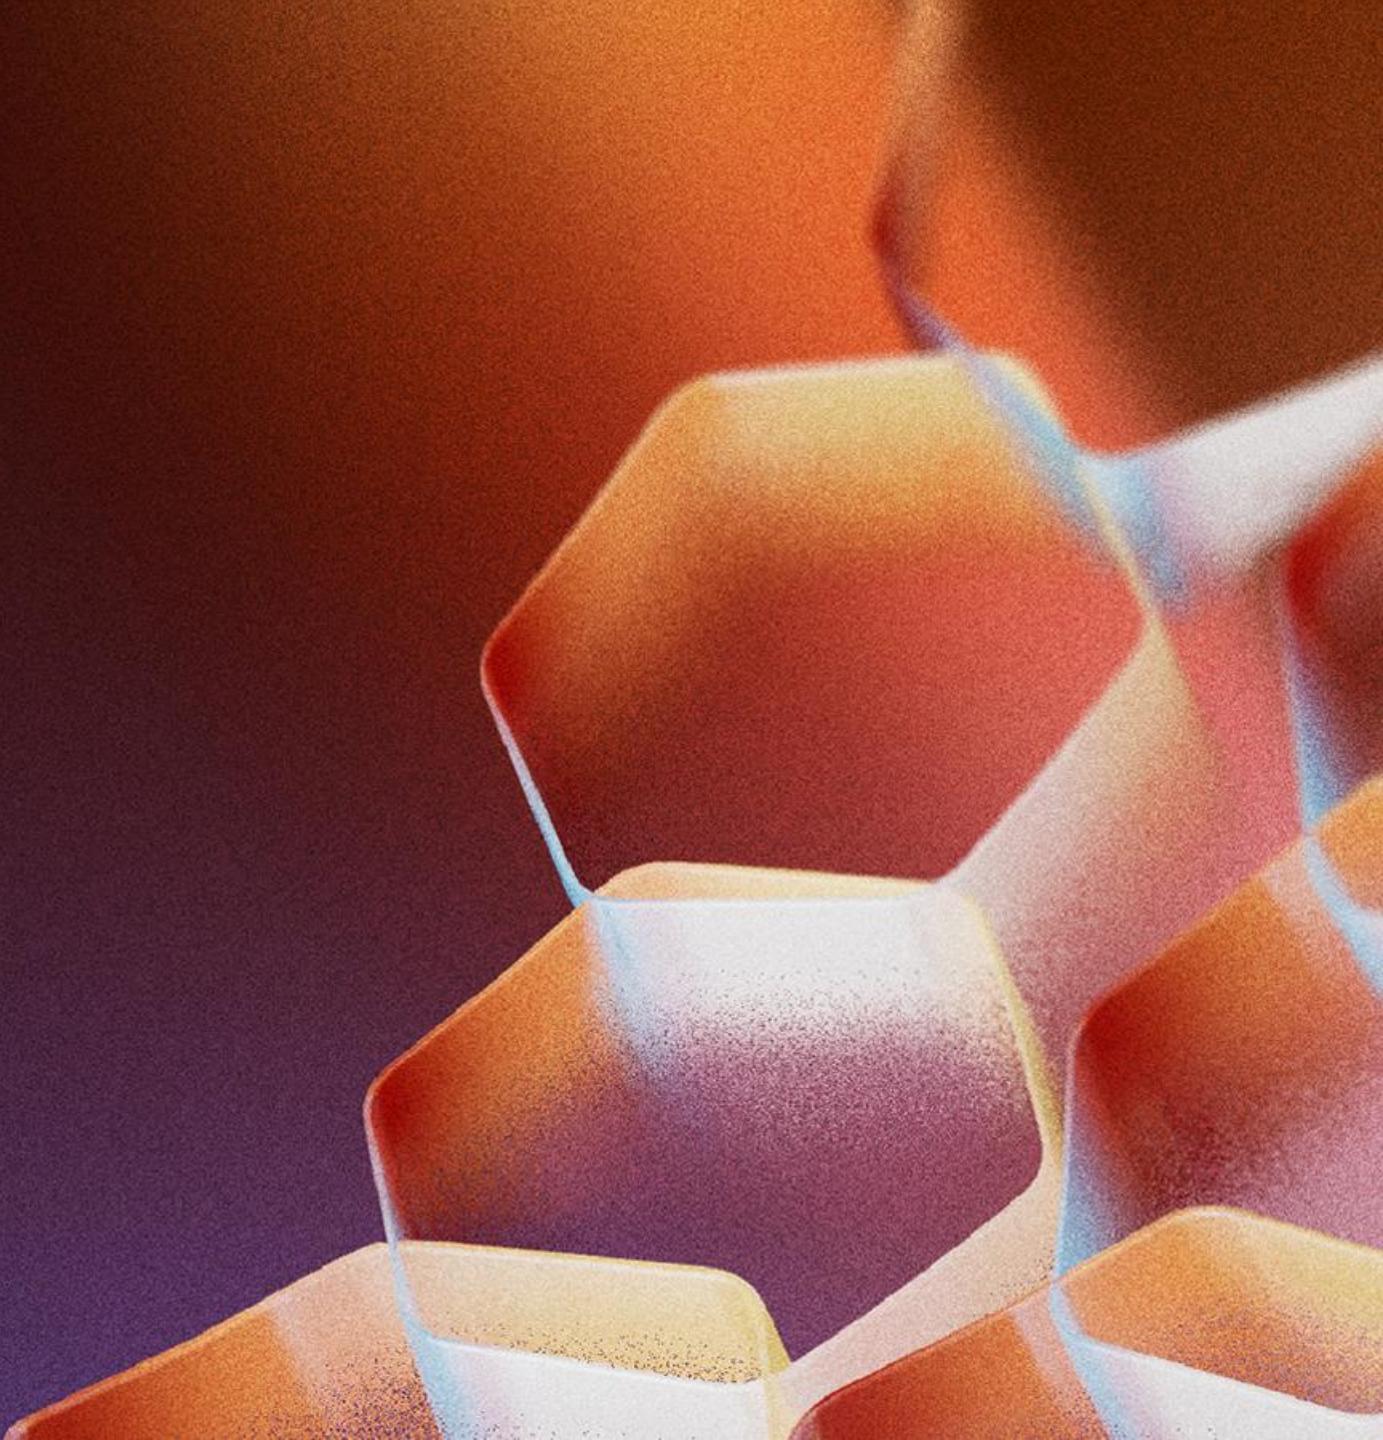
# Change management process

**Remember to accept all your shifts and declining a shift calendar hold does NOT remove you from the shift!**

- Final shift calendar holds were sent on May 13. If you do not see any shifts on your calendar, please reach out to [ESCExpert@microsoft.com](mailto:ESCExpert@microsoft.com).
- Please submit any shift change requests to your Expert Recruit Lead. Change requests will be reviewed by the planning team, and approvals will be determined in collaboration with Expert Recruit Leads.
- As we are currently T-1, change requests without a critical business justification will be denied.
- If you need to step away for any reason, please work directly with your topic lead to ensure coverage while you are away.



# Logistics



# Badge Pickup

- Badge Pickup is located in the Atrium on Level 4
  - 705 Pike Arch Seattle Convention Center.
- Pick up your badge on May 18<sup>th</sup> to avoid lines!
  - Sun, May 18: 9:00 AM – 6:00 PM
  - Mon, May 19: 7:00 AM – 7:30 PM
  - Tues, May 20: 7:00 AM – 6:30 PM
  - Wed, May 21: 7:30 AM – 6:00 PM
- You must present government issued photo identification to receive your Microsoft Build Badge.
- The Hub will be open for early access for Experts on Sunday, May 18 from 9:00 AM – 5:00 PM.

## Speaker & Expert Hoodies

- Approved in-person Speakers and Experts registered by May 1 will have the option to order an event hoodie of their preferred style and size, available for pick-up onsite.
- Speakers and Experts registering after May 1 will be eligible to receive an event hoodie on a first come, first served basis.
- Hoodies will be available for pick-up in the speaker workroom in 602-603 on Level 4 of 705 Pike Arch Seattle Convention Center, starting May 18.



# R&Rs for Expert Meet-up Staff

## Early Access



- Expert early access hours are 11 AM – 5 PM on 5/18. On event days, Experts will have access 30 mins before Hub open.

- Experts can view the Keynote livestream in the demo theaters and F&B space on 5/19 & 5/20.

## Shift Check-in



- Check-in 30-mins before the start of each shift at the Expert Shift Check-in desk in Level 4 Atrium.

- Keep personal items to a minimum and place them in the storage cabinet during your shift.

## Expert Shift



- Remain at your assigned station during the entirety of your shift.

- Do not leave your station unattended, coordinate break coverage with your ERL or other experts.

## Conversation



- Engage with attendees and answer questions related to your topic of expertise.

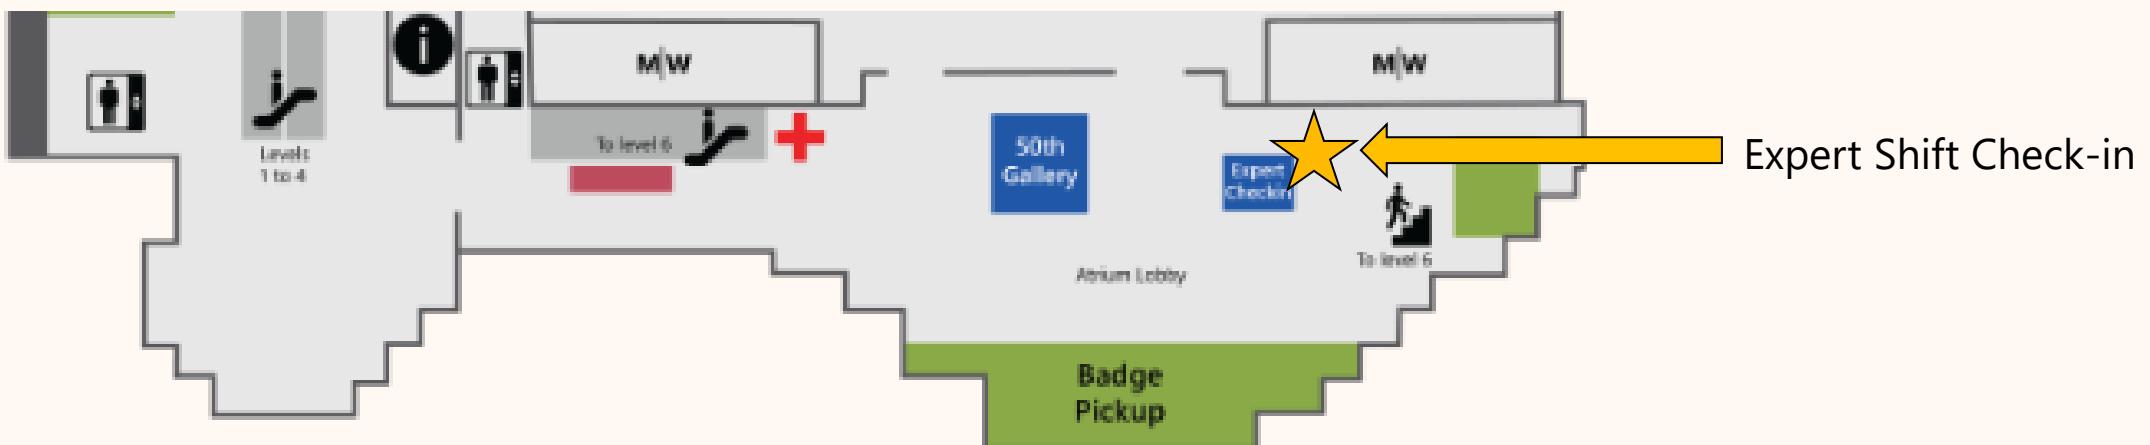
- Be familiar with sessions and partners related to your topic of expertise and make recommendations.

## Content



- If you plan to demo, you may use any content that your ERL has pre-loaded on the station device.

- You may also bring materials on a USB stick or use your own device (HDMI adaptor required).

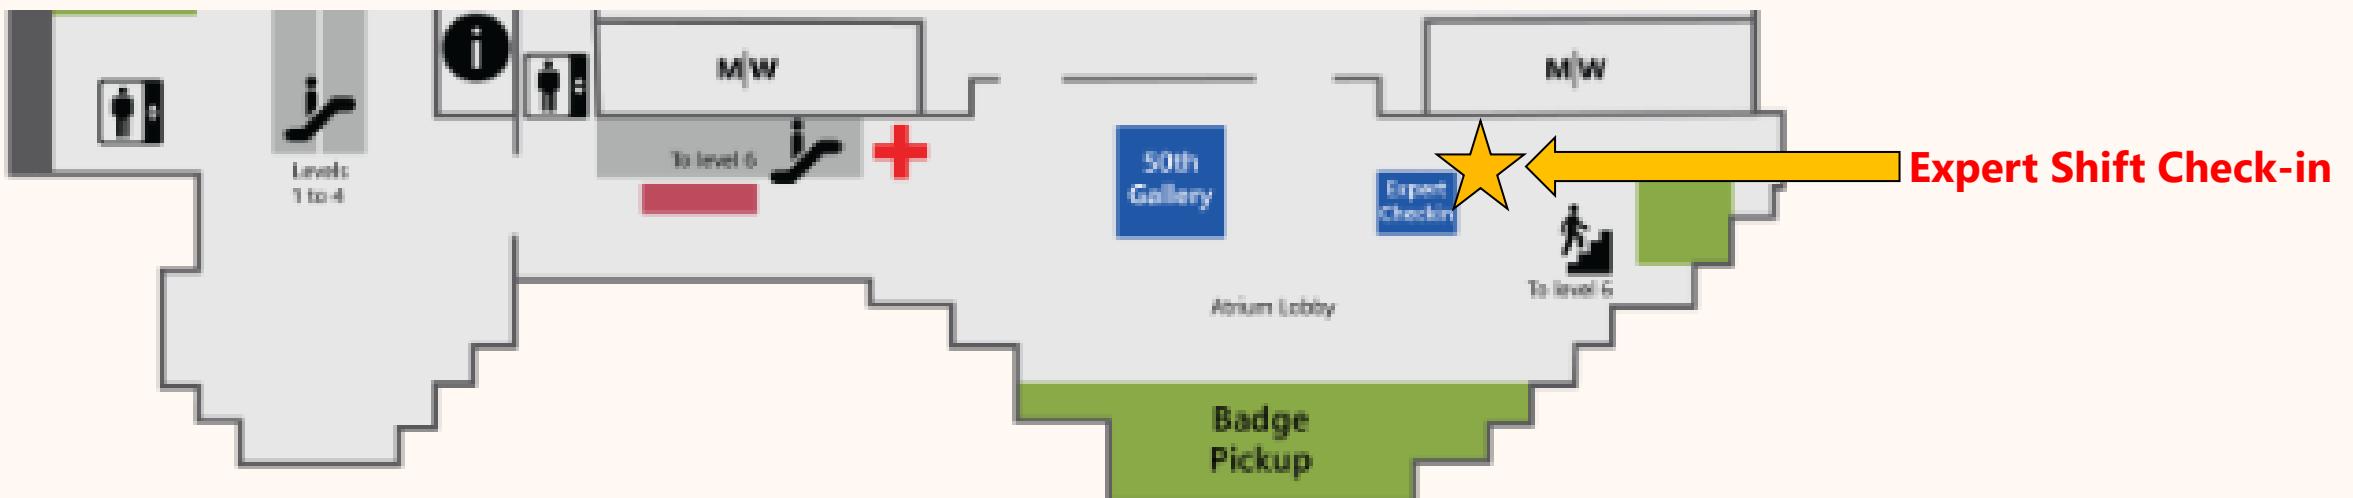


# Expert shift check-in

## IMPORTANT!

**Check-in for each shift** – Remember to check-in for EACH shift at the Expert shift check-in desk.

- Desk will be located in the Atrium on Level 4, next to Badge Pickup and in front of the Hub entrance.
- This desk will be staffed by EventForce and supported by a member of the EMU events team.
- Please check-in at the start of EACH scheduled shift, not just your first shift!
- If you forgot to check-in, notify your ERL or stop by the check-in desk during/after your shift.



# Media Engagement

Tips below are from PR's conference show floor guidance doc:

**Stay within your designated scope** - Once you've established your area of expertise, stay there. You have no obligation and should not be commenting on anything other than your area of focus. If asked broader questions, it's perfectly ok to explain that you're not the right person to address these questions.

**Remember "The New York Times" rule** - Before responding to any question, think about how it would look as a headline or quote in The New York Times. If it's something you (or your manager) would not want to see as a headline or a quote in The Times, don't say it.

**"No comment" is never the right answer** - If ever pressed to share confidential information, rely on a polite, but firm statement, "That's not something I can share," or "I'm not the right person to speak about that," for example. Kindly offer to put them in touch with your lead to help with questions or guidance.

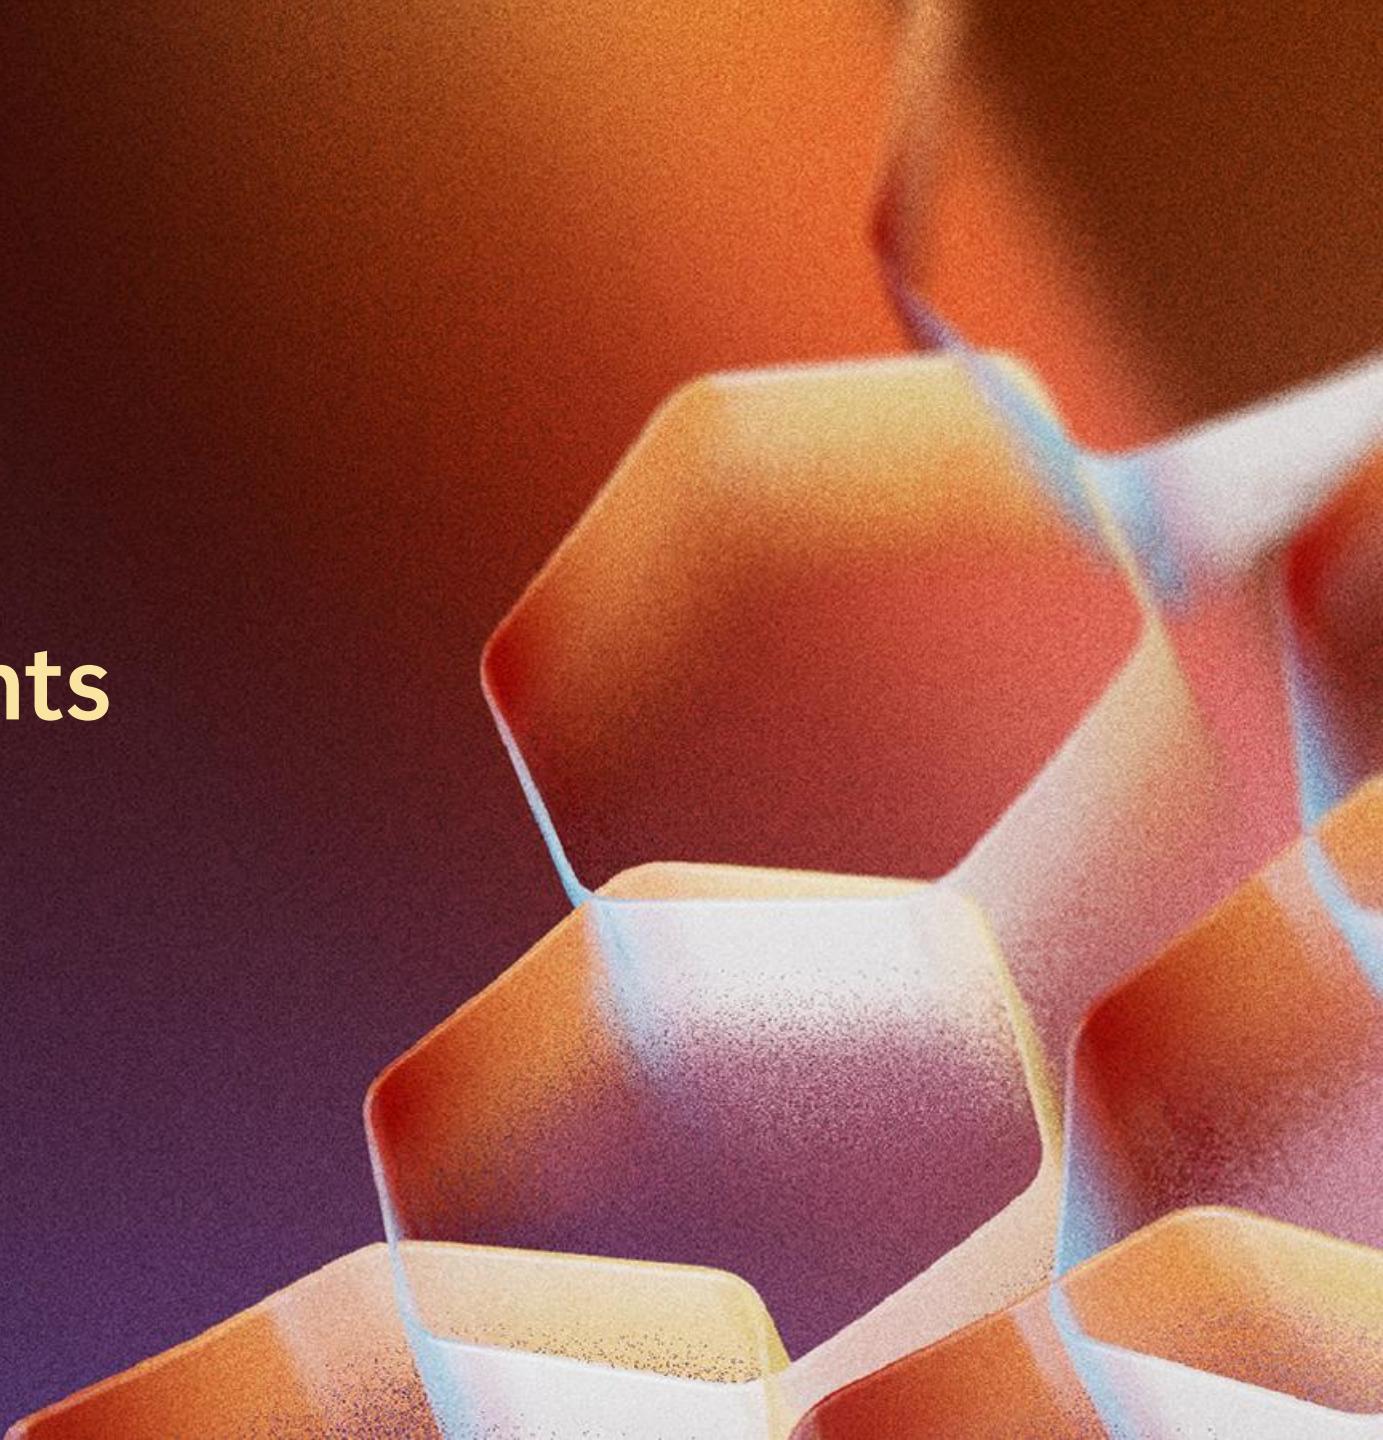
**Know your communications contact** - If approached by a reporter with broader questions or company commentary beyond your demo/area of focus, politely put them in touch with your lead. The interview opportunity will be evaluated, and they will be connected with an appropriate corporate spokesperson, as needed.



# Product announcements

Jenn Cockrell

[AI Foundry Build 25 Sessions .xlsx](#)



# Azure AI Foundry—Microsoft Build 2025

## Models & fine-tuning

- New Foundry Models & partnerships, including unannounced models & purchase flexibility
- New model leaderboard, router, and video playground
- Fine-tuning support for additional models and support for reinforcement fine-tuning
- Foundry Local for on-device deployment (including Windows & Mac)
- New models and tools from Microsoft Research added to Foundry Labs

## Agents, knowledge, & tools

- Foundry Agent Service general availability with agent IDs, agent catalog, BYO LLM from any cloud, a single SDK for Semantic Kernel and AutoGen, grounding and agentic RAG
- Support for A2A and MCP interoperability. LangChain and CrewAI tooling coming soon
- New Azure AI services for Content Understanding and Voice API

## Observability and Trustworthy AI

- Enhanced security and safety capabilities in Foundry
- Foundry Observability enhancements

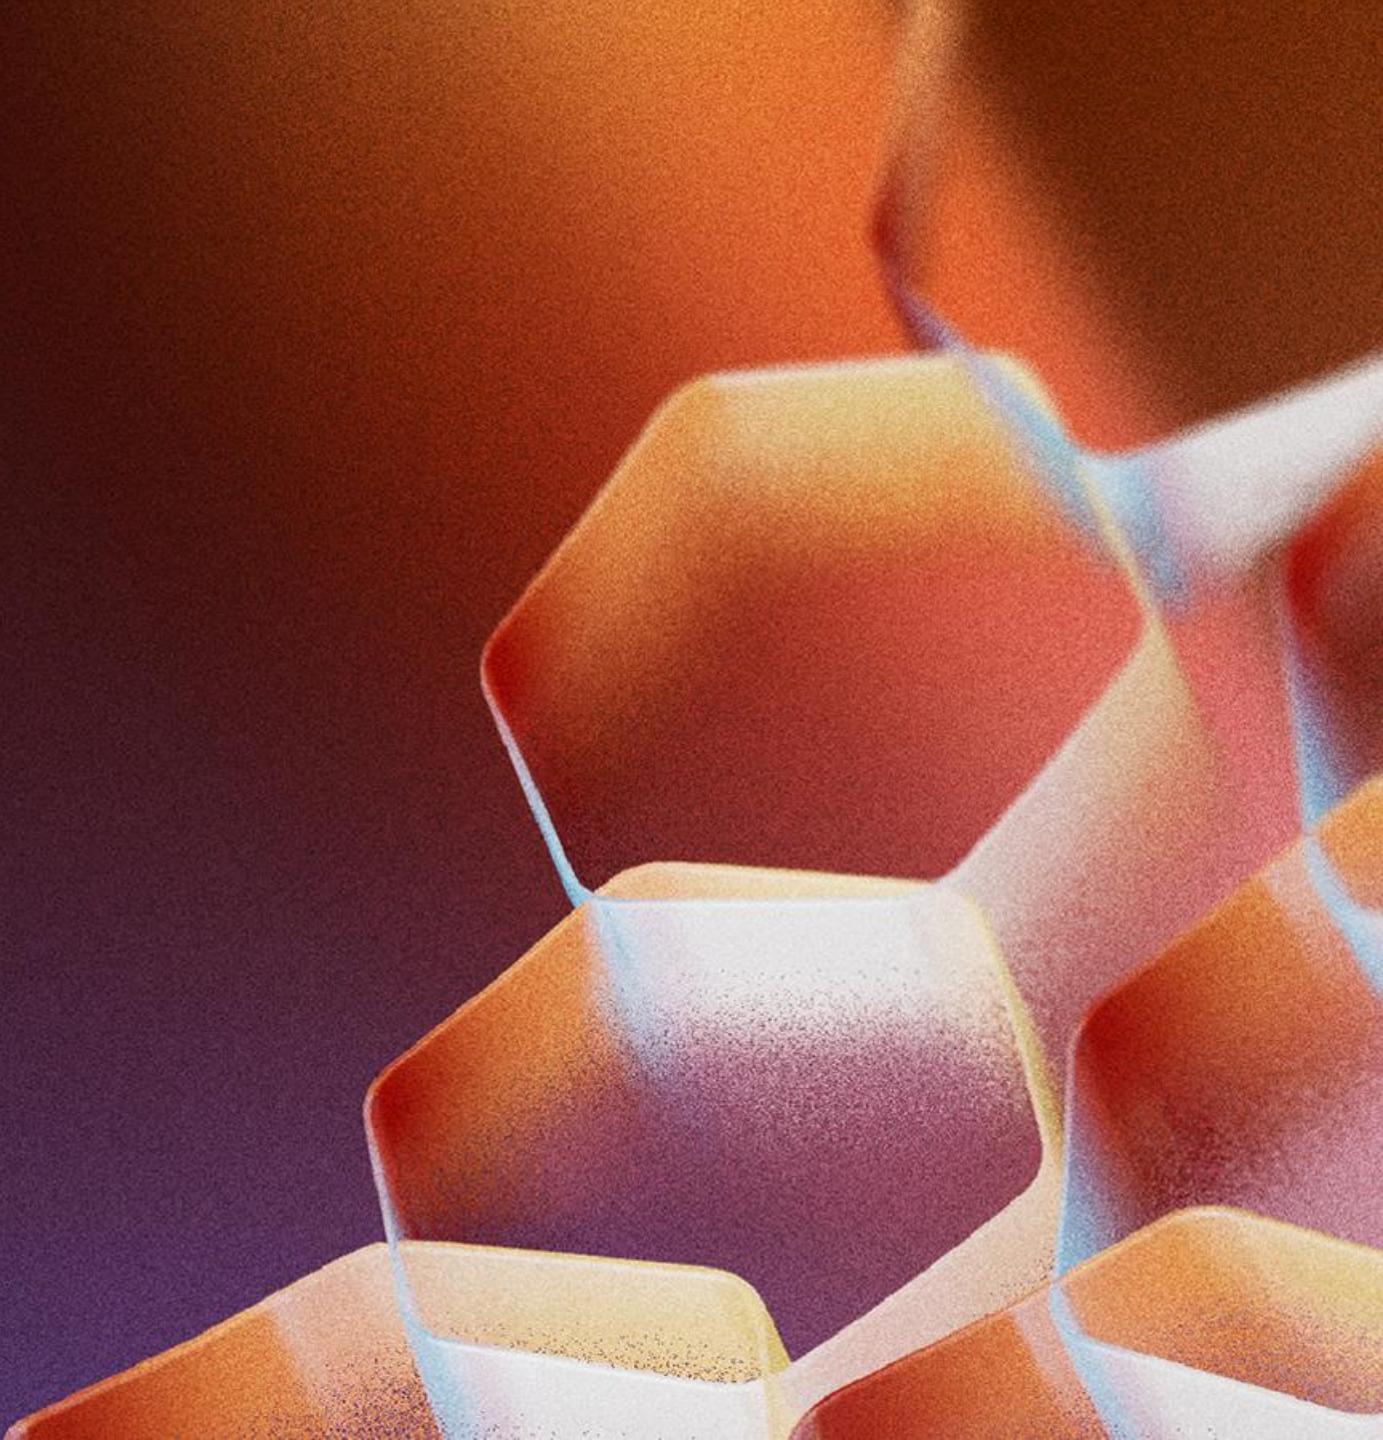
## Developer workflow improvements

- Azure AI Foundry API general availability and Azure AI Foundry projects client library
- 10 new quickstart templates
- New Azure AI Foundry capabilities in Copilot Studio and GitHub Codespaces
- Azure AI Foundry extension for Visual Studio Code (announced 4/4)



# Resources

Nick Brady & Mark Baiza



# Azure AI Foundry content repository

The screenshot shows a GitHub repository page for 'build-2025-demos'. The repository is private and has 1 branch and 0 tags. It contains two commits: one from 'microsoft-github-operations[bot]' and another from '.gitignore'. The '.gitignore' file was committed 4 days ago. The repository has 1 star, 2 forks, and 1 watch. The 'About' section describes it as 'AI demos at the Expert Meetup booths showcasing models, agents, and tools.' The 'Code of conduct' section links to the Microsoft Open Source Code of Conduct. The 'Releases' and 'Packages' sections both indicate 'No releases published' and 'No packages published' respectively.

azure-ai-foundry / build-2025-demos

Type  to search

Code Issues Pull requests Actions Projects Security Insights

Watch 1 Fork 2 Star 1

build-2025-demos Private

main 1 Branch 0 Tags Go to file Add file Code

microsoft-github-operations[bot] Initial commit b596e60 · 4 days ago 1 Commit

.gitignore Initial commit 4 days ago

Code of conduct

## Microsoft Open Source Code of Conduct

This project has adopted the [Microsoft Open Source Code of Conduct](#).

Resources:

- [Microsoft Open Source Code of Conduct](#)
- [Microsoft Code of Conduct FAQ](#)
- Contact [opencode@microsoft.com](mailto:opencode@microsoft.com) with questions or concerns

About

AI demos at the Expert Meetup booths showcasing models, agents, and tools.

Code of conduct

Activity

Custom properties

1 star

1 watching

2 forks

Releases

No releases published

Packages

No packages published

Microsoft Open Source Docs Portal GitHub @ MSFT Support Options Data Protection Notice

© 2025 GitHub, Inc. Terms Privacy Security Status Docs Contact Manage cookies Do not share my personal information

# Azure AI Foundry | Feedback collection process

<https://aka.ms/AI/Feedback>

 **Product Feedback and Ideas**

Provide information about using the product. This could include areas such as user experience, functionality, features, bugs and issues. As well as needs the platform may not provide today

 1



**Agents not cleaned up Ai foundry – Duplicate agent instances created on each run in Multi-Agent Accelerator**

[ai-foundry-general](#) [ai-agents](#) [semantic-kernel](#)

 UtkarshMishra-Microsoft asked last week in [Product Feedback and Ideas](#)

 3

- Answered

# Azure AI Foundry | Gartner Analyst Peer Insights

## Program Objective

*The program objective is to increase the quantity and quality of product review submissions for Azure AI Foundry on Gartner to **improve our standing as a leader in tier 1 analyst reports**, in support of **establishing Azure as the #1 cloud for AI**.*

## Our Purpose

- **Help empower prospects and decision-makers with peer-to-peer guidance** to make informed purchase decisions during prospect research phase
- **Improve Azure AI Foundry's standing as a leader in tier 1 analyst reports** for Gartner, IDC and Forrester
- **Elevate Azure AI Foundry's brand reputation across the industry** in support of establishing Azure as the #1 cloud for AI

## Our Strategy

- **Have customers share** how Azure AI Foundry has helped them solve business critical problems and created value for their business via online reviews
- **Streamline process** for gathering reviews
- **Help PMs and PMMs identify opportunities** and weaknesses in products and customer experience
- **Help sales understand the customer voice** and close deals with peer endorsements
- **Use proof points and quotes in marketing materials**



## Leave a product review, get \$25!

- Share your insights using Azure AI Foundry. Help other developers and become a thought leader.
- Choose a **\$25 gift card** as a “thank you” for your time.
- Any technical roles, from VPs to developers, can leave a review. Must be a real customer of Azure AI Foundry.

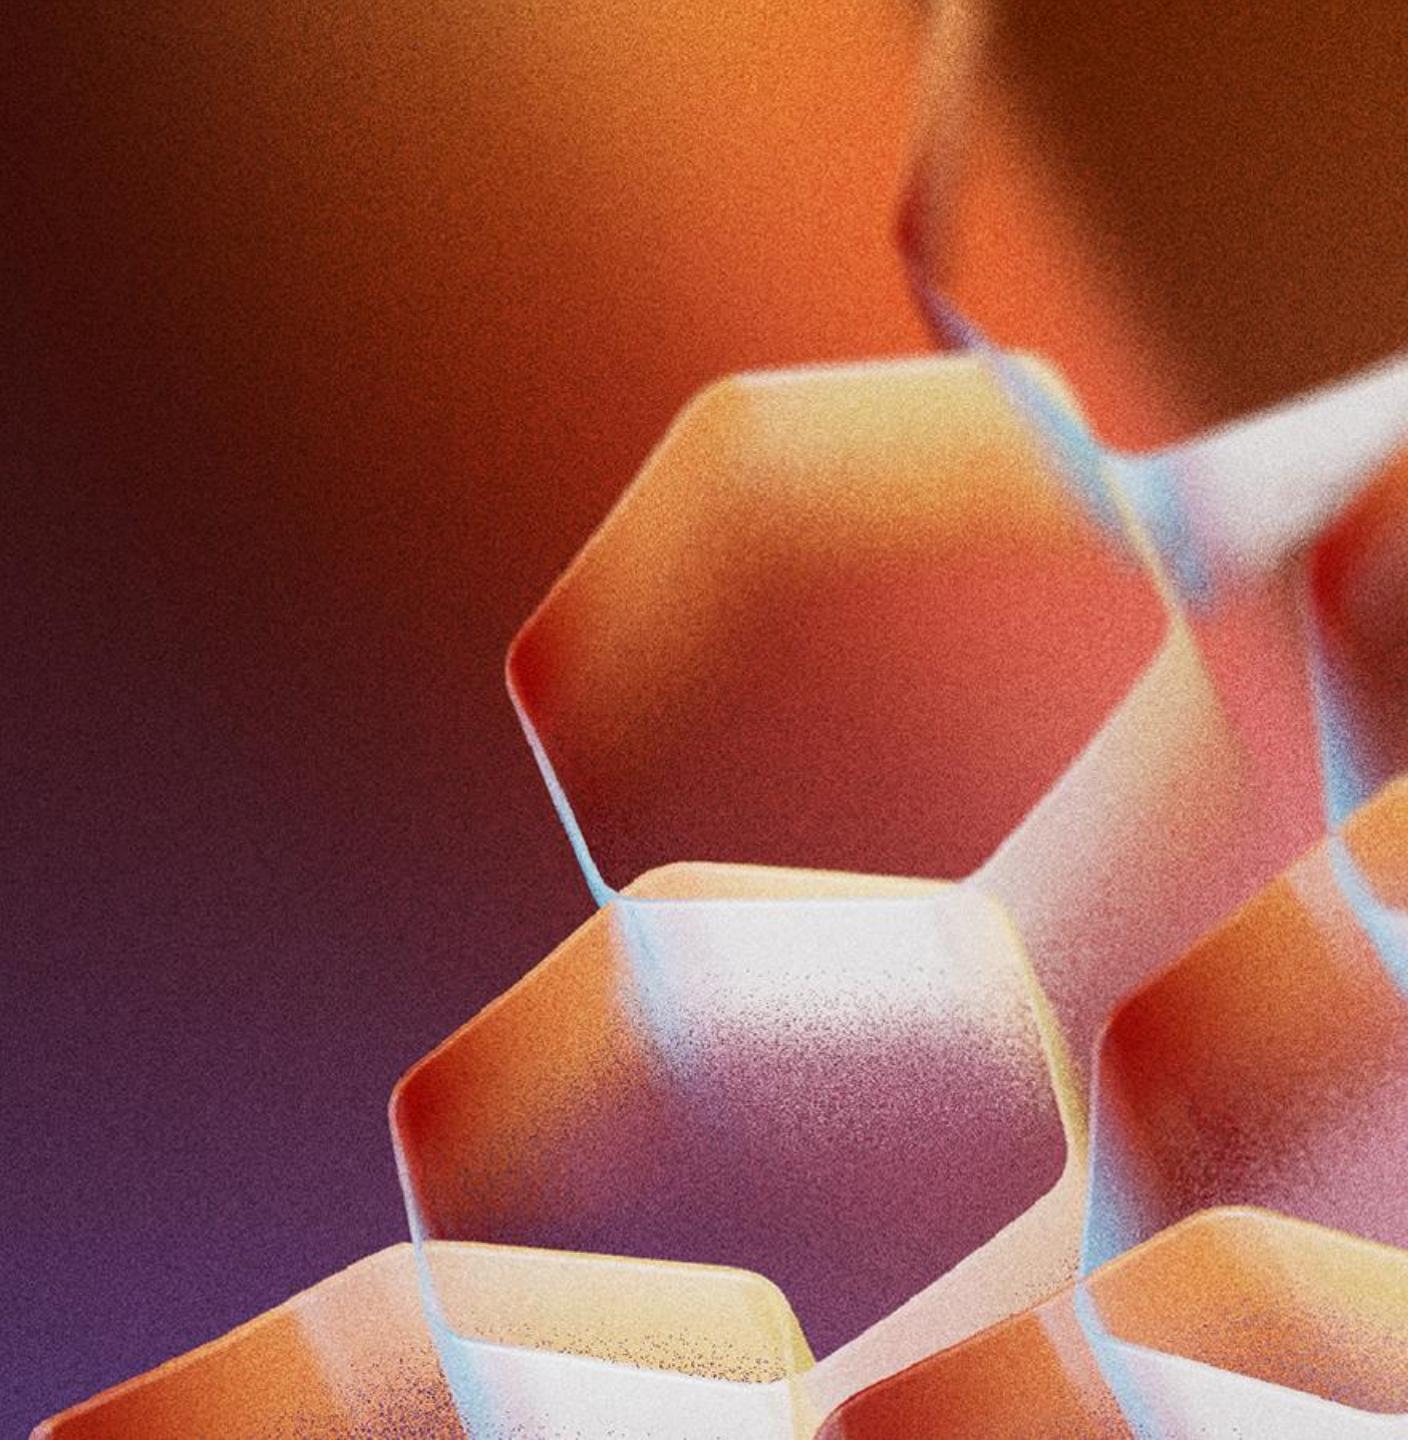


<https://aka.ms/leaveareview>



# AI Expert Leads

Laura Wolf



---

# Highlights of the AI Expert Meetup

- Designed to foster alignment and knowledge sharing
- Engages developer audience during Build
- Brings together technical leads, marketing, and Anand's team
- Offers personalized 1:1 or small group conversations
- Features demo-focused planned content delivered by Microsoft experts and MVP SMEs
- Addresses technical questions and provides business solutions
- Targets key segments: back-end, front-end, and full-stack developers

# AI Expert Meet-up booths

## 10. Azure AI Tools

Elevate applications with speech, language, and translation

## 11. Azure AI Foundry Observability

Govern workflows using robust tracing, evaluation, and observability

## 12. Azure AI Foundry Content Safety

Detect and filter harmful content with advanced safeguards

## 13. AI Skilling

Advance your AI expertise

## 14. Azure AI Document Intelligence

Transform your content

## 15. Foundry Local

Run AI models and agents on-device

## 16. Azure AI Search Agentic Retrieval

Precision in information retrieval

## 17. Model Customization

Tailor models to your needs with fine tuning

## 18. Azure OpenAI in Foundry Models

Models designed for complex problem-solving

## 19. Azure AI Foundry Models

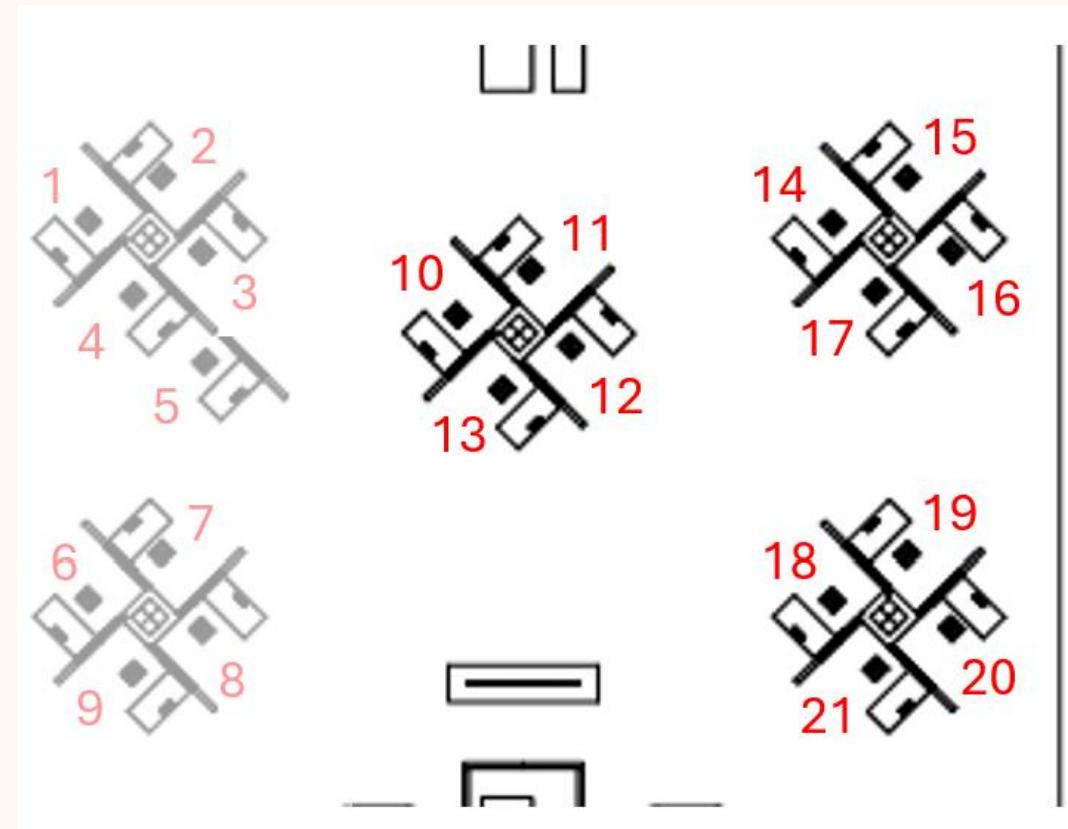
Discover the best model for your use case

## 20. Azure AI Foundry Agent Service

Orchestrate agentic workflows

## 21. Azure AI Foundry

Create generative AI applications and agents

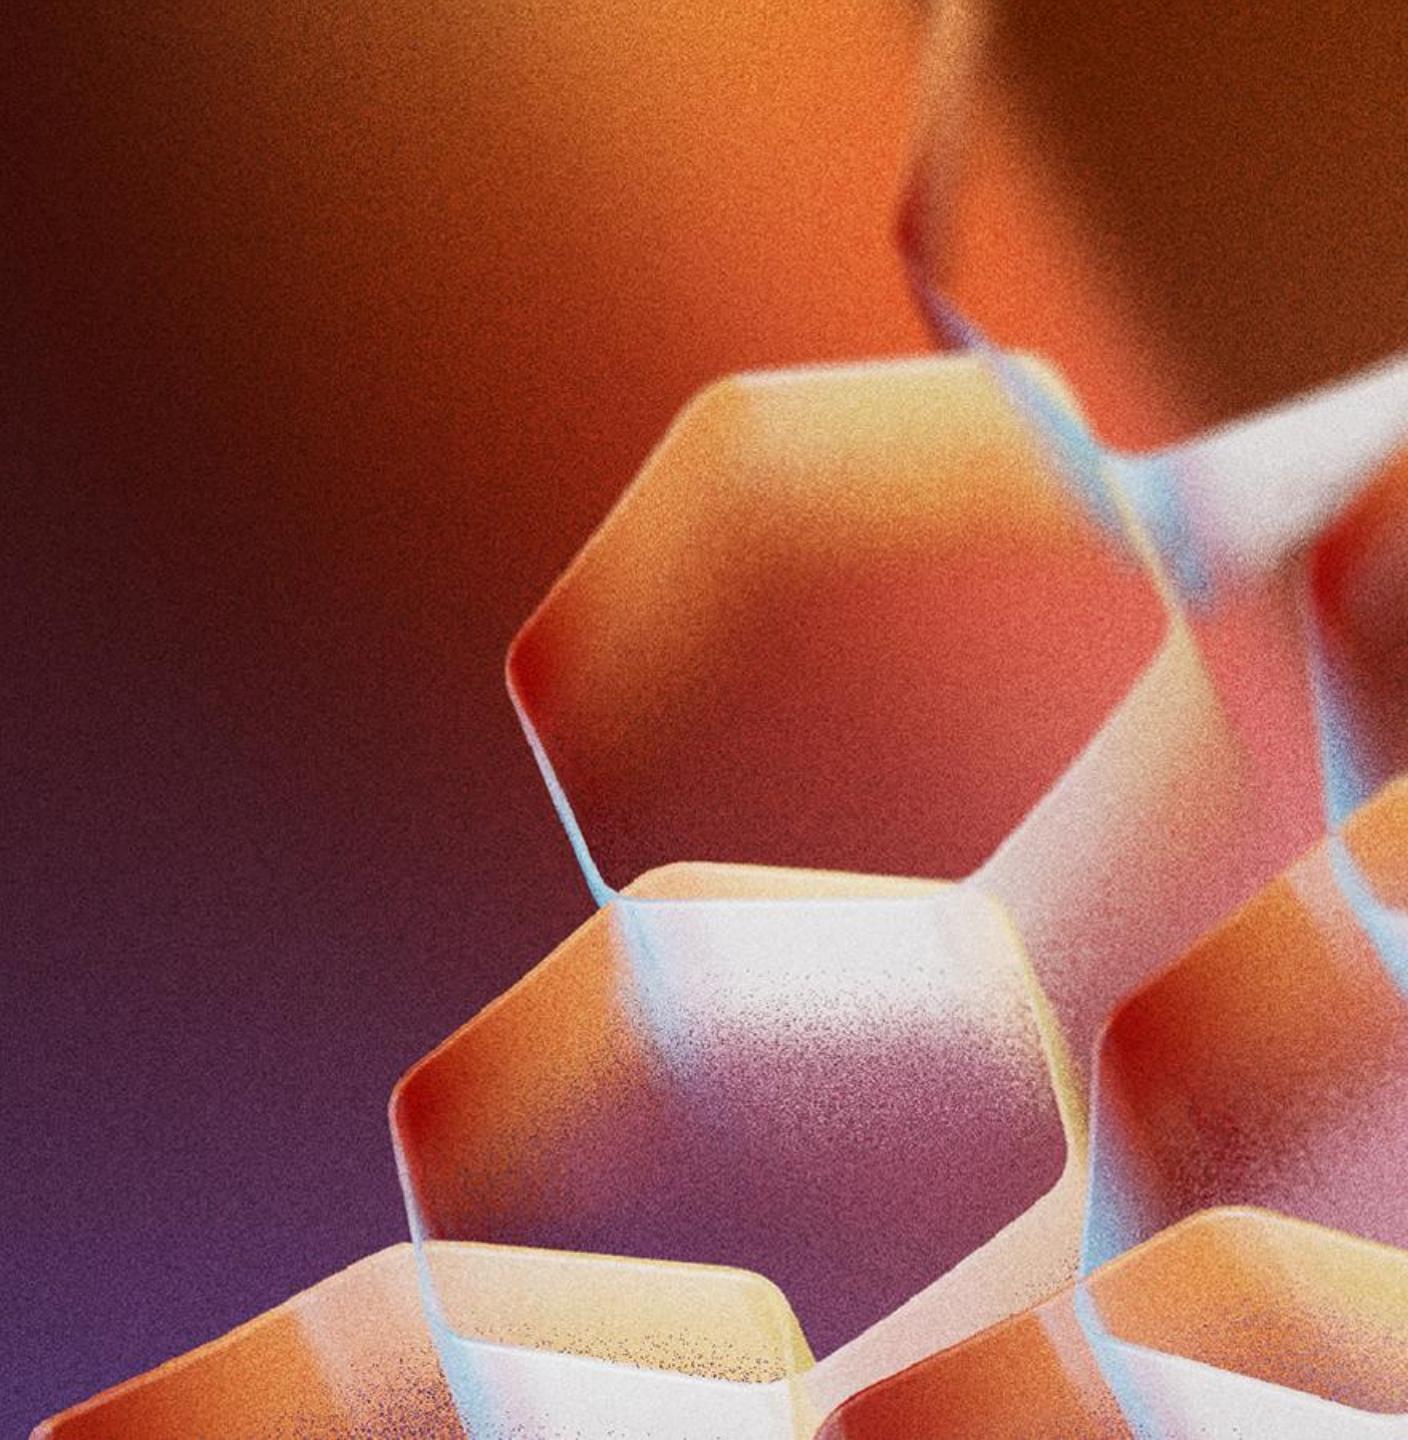


# Azure AI Foundry | Expert Leads

<b>Quad</b>	<b>Station Placement</b>	<b>Azure AI Foundry Station topics</b>	<b>Lead</b>	<b>PMM Lead</b>
1	21	Azure AI Foundry	Nick Brady	Courtney Brewer
	19	Azure AI Foundry Models	Sara Kandil	Sharmi Chockalingam
	18	Azure OpenAI in Foundry Models	Vanessa Diaz	Andy Beatman
	20	Azure AI Foundry Agent Service	Pavan Li	Takuto Higuchi
2	16	Azure AI Search Agentic Retrieval	Mike Carter	Allison Sparrow
	17	Model Customization	Nandini Muralidharan	Chhavi Nijhawan
	15	Foundry Local	Jhoanna Machado	Richard Tso
	14	Azure AI Document Intelligence	Aditi Maheshwari	Ronak Chokshi
3	10	Azure AI Tools	JJ Kim	Takuto Higuchi
	12	Azure AI Foundry Content Safety	Ying Chang	Katelyn Rothney
	11	Azure AI Foundry Observability	Sebastian Kohlmeier	Tara Grumm
	13	AI Skilling	Lanna Teh	n/a



# Focus area deep dive





# Azure AI Foundry

Nick Brady

# The Agent Factory | BRK155 | Monday, 11:45a

The image displays three side-by-side screenshots of the Azure AI Foundry platform:

- Models:** Shows the "Model catalog" page with sections for "Find the right model to build your custom AI solution" and "Model leaderboards". It lists models like Mistral Medium 3, New Phi reasoning models, Introducing gpt-image-1, MAI-DS-R1 reasoning model, and Unlock with o3. Below are "Model leaderboards" for Quality, Safety, Cost, and Throughput.
- Agents:** Shows the "Agents" management page. It lists four agents: "Agent185" (Retriever, gpt-4), "Agent186" (Unifier, gpt-4), "Agent187" (Planner, gpt-4), and "Agent188" (Planner, gpt-4). The "Setup" section shows deployment details for "Agent185" using "Grounding with Bing S..." and "gpt-4s".
- Tools:** Shows the "Agents playground" page for a specific thread. It displays a list of tools: "bing\_grounding", "test\_LLM", "test\_JAIS", "test\_JAIS", "test\_JAIS", "test\_JAIS", and "message\_create". The "Input" field contains the URL "https://api.bing.microsoft.com/v7.0/search/q=reggaean+site:rejigl.com?".

Models

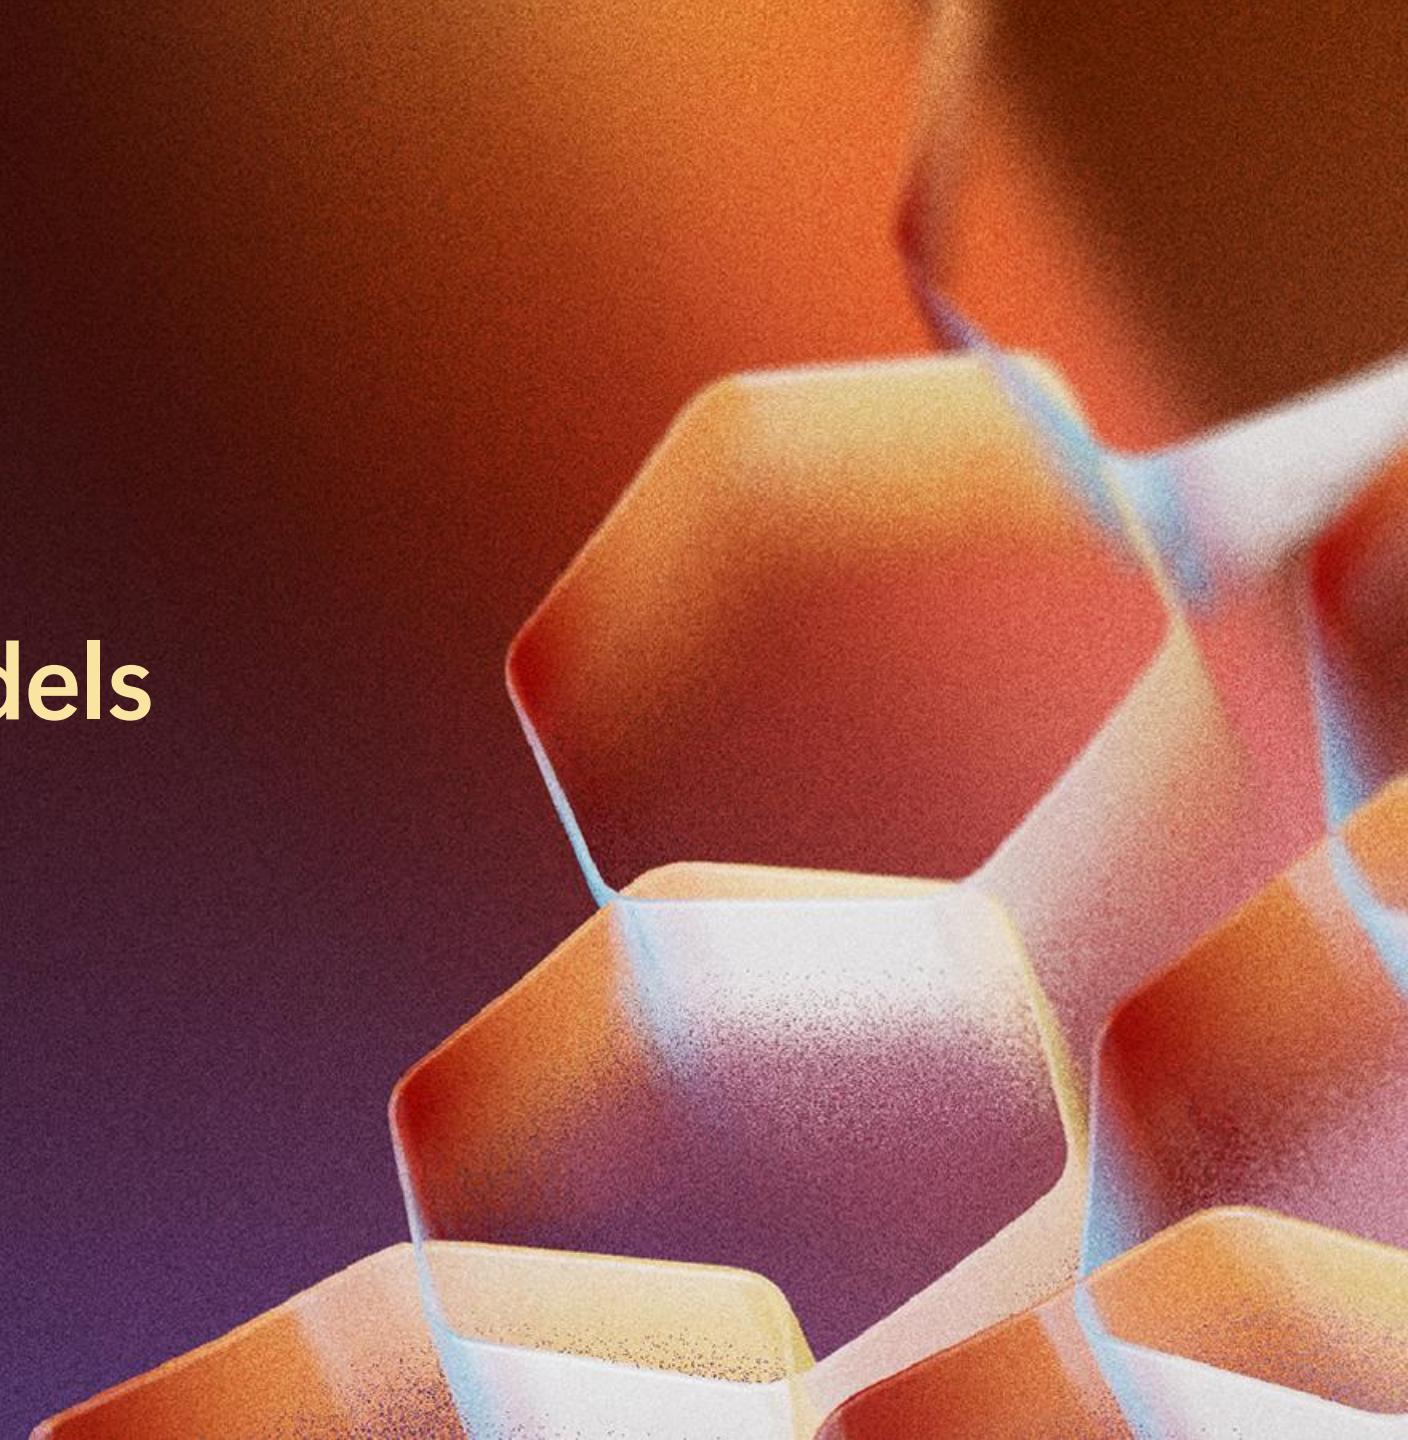
Agents

Tools



# Azure AI Foundry Models

Sara Kandil



# Azure AI Foundry Models



Sara Kandil

## Brief introduction

- 🚀 Part of the Applied AI Engineering team
- 🤝 Focused on empowering Microsoft Partners to accelerate adoption of Azure AI Foundry
- 🎉 Over 5 years at Microsoft
- 🎤 First time attending Build — super excited to be here!

## Topics Overview (WIP)

- Foundry Models Fungible PTU
- Nvidia NIM integration and MaaP monetization
- Foundry Models MCP server
- + More!

## Current Content Plan (WIP) & Work in Progress

- Introduced myself to Manoj and Sharmila
- Received booth topics overview from Manoj
- Contact all experts to confirm shifts, introductions, and on-time attendance
- [Thursday] Finalize content and demos
  - [OQ] Will internal environments be accessible for live demos?
- [Thursday] Scheduled a recorded prep meeting with all experts and product team, where they will walk through topics and demos
- [TBD] Upload resources to Github Repo



# Azure OpenAI in Foundry Models

Vanessa Diaz

# Azure OpenAI in Foundry Models

Tier 1 & 2 Announcements	Content (avail on GH TBA)	Sessions
<ul style="list-style-type: none"><li>• Models</li><li>• Feature enhancements</li></ul>	<ul style="list-style-type: none"><li>• Announcement slides</li><li>• Demos</li><li>• Videos</li></ul>	<ul style="list-style-type: none"><li>• <a href="#"><u>BRK 173 Advanced Azure OpenAI Innovation in Azure AI Foundry</u></a></li><li>• <a href="#"><u>BRK170 Revolutionizing AI Apps with Multimodal Models in Azure AI Foundry</u></a></li><li>• <a href="#"><u>BRK150 Fine-tuning and distillation with Azure AI Foundry</u></a></li><li>• <a href="#"><u>BRK178 Optimize your GenAI applications at scale in Azure AI Foundry</u></a></li></ul>



# Azure AI Foundry Agent Service

Pavan Li

# Azure AI Foundry Agent Service

Tier 1 & 2 Announcements	Content (avail on GH TBA)	Sessions
<ul style="list-style-type: none"><li>Foundry Agent Service general availability</li><li>Connected agents and multi-agent workflows (in preview), integration with Semantic Kernel</li><li>Knowledge tools expansion with SharePoint and 3P partner tools</li><li>Agent Catalog that hosts a library of agent code samples</li><li>Support for A2A and MCP interoperability. LangChain and CrewAI tooling coming soon</li><li>AgentOps tooling including evaluation, tracing, and monitoring</li><li>BYO-thread storage with Cosmos DB</li></ul>	<ul style="list-style-type: none"><li>Announcement slides</li><li>Demos</li><li>Videos</li><li>Blogs</li></ul>	<b>BRK155</b> <a href="#">Azure AI Foundry: The Agent Factory</a>
		<b>BRK149</b> <a href="#">Azure AI Agent Service: Transforming workflows with Azure AI Foundry</a>
		<b>BRK148:</b> <a href="#">Building the digital workforce: Multi-agent apps with Azure AI Foundry</a>
		<b>BRK152:</b> <a href="#">Best practices for building agentic apps with Azure AI Foundry</a>



# Azure AI Foundry

## AI Agent Ecosystem

### Authoring Tools

Visual Studio

GitHub

### Knowledge



Microsoft Fabric



Azure Cosmos DB



SharePoint



Microsoft Graph



Azure AI Search



Microsoft Bing

### Enterprise Trust

VNet deployments

OBO Auth

BYO-resources

Managed Identities

### Agent catalog

*Ready-made agents to kickstart your agent workforce*

### Agent frameworks

*Client SDKs for enterprise and production agentic systems*

Agent tools

Orchestration

Channels

### Foundry Agent Service

*Deploy and manage agents with fully-managed runtime*

Multi-agent workflows

Built-in threads

Long-term memory

Evaluation

Tracing & Monitoring

Governance + Safety

### Foundry Models

#### Azure Direct Models

OpenAI | Llama | Grok | Flux | Mistral

- Phi
- Stability.ai
- Cohere
- Paige
- Hugging Face
- Industry Models

#### Azure Ecosystem Models



### Foundry Labs

Muse

Magnetic-One

Aurora

OmniParser

### Open Ecosystem

#### Protocols

A2A | MCP | OpenAPI

#### Connectors

Google VertexAI

Amazon Bedrock

CrewAI

LangChain

ElasticSearch

Pinecone

Logic Apps

Azure Functions



# Azure AI Foundry Agent Service

Generally available

Multi-agent  
workflows

Open and  
interoperable tooling

Trust and enterprise  
security



# Azure AI Search Agentic Retrieval

Mike Carter

# Azure AI Search Agentic Retrieval

Tier 2 Announcements	Event Coverage	Release type	Content
<b>Agentic retrieval: declarative, multi-turn query engine</b> <ul style="list-style-type: none"><li>Query planning with conversation history</li><li>Extractive answers</li><li>Token-based billing model</li><li>Demo: Agentic RAG with Agent Service contextual memory</li></ul>	<ul style="list-style-type: none"><li>Keynote</li><li><a href="#">BRK141</a></li><li><a href="#">BRK142</a></li></ul>	Public preview	<ul style="list-style-type: none"><li>Session slides</li><li>Demo</li><li>2 blog posts</li><li>Docs</li></ul>
<b>Multimodal ETL</b> <ul style="list-style-type: none"><li>GenAI prompt skill: Enrich data on ingestion with chat completion &amp; <a href="#">model inference</a></li><li>Demo: Multimodal ingestion in the Azure portal</li></ul>	<ul style="list-style-type: none"><li><a href="#">BRK141</a></li></ul>	Public preview	<ul style="list-style-type: none"><li>Session slides</li><li>Demo/Sample</li><li>2 blog posts</li><li>Docs</li></ul>
Tier 3+ Announcements			
<b>Enterprise security in Azure AI Search</b> <ul style="list-style-type: none"><li>Doc level security with ADLSGen2 ACL support</li><li>Purview security labels (private preview)</li></ul>	<ul style="list-style-type: none"><li><a href="#">BRK141</a></li></ul>	Public preview	<ul style="list-style-type: none"><li>Session slides</li><li>Blog post</li></ul>
<ul style="list-style-type: none"><li>Change your pricing tier (SKU migration) &amp; Service Upgrade</li></ul>		GA	<ul style="list-style-type: none"><li>Blog post</li><li>Docs</li></ul>
<b>Azure AI Search as a knowledge source</b> <ul style="list-style-type: none"><li>Copilot Studio</li><li>AI Agent Service</li><li>Knowledge in Foundry</li></ul>	<ul style="list-style-type: none"><li><a href="#">BRK149</a></li><li><a href="#">BRK155</a></li></ul>		<ul style="list-style-type: none"><li>Session slides</li></ul>

## Azure AI Search Sessions:

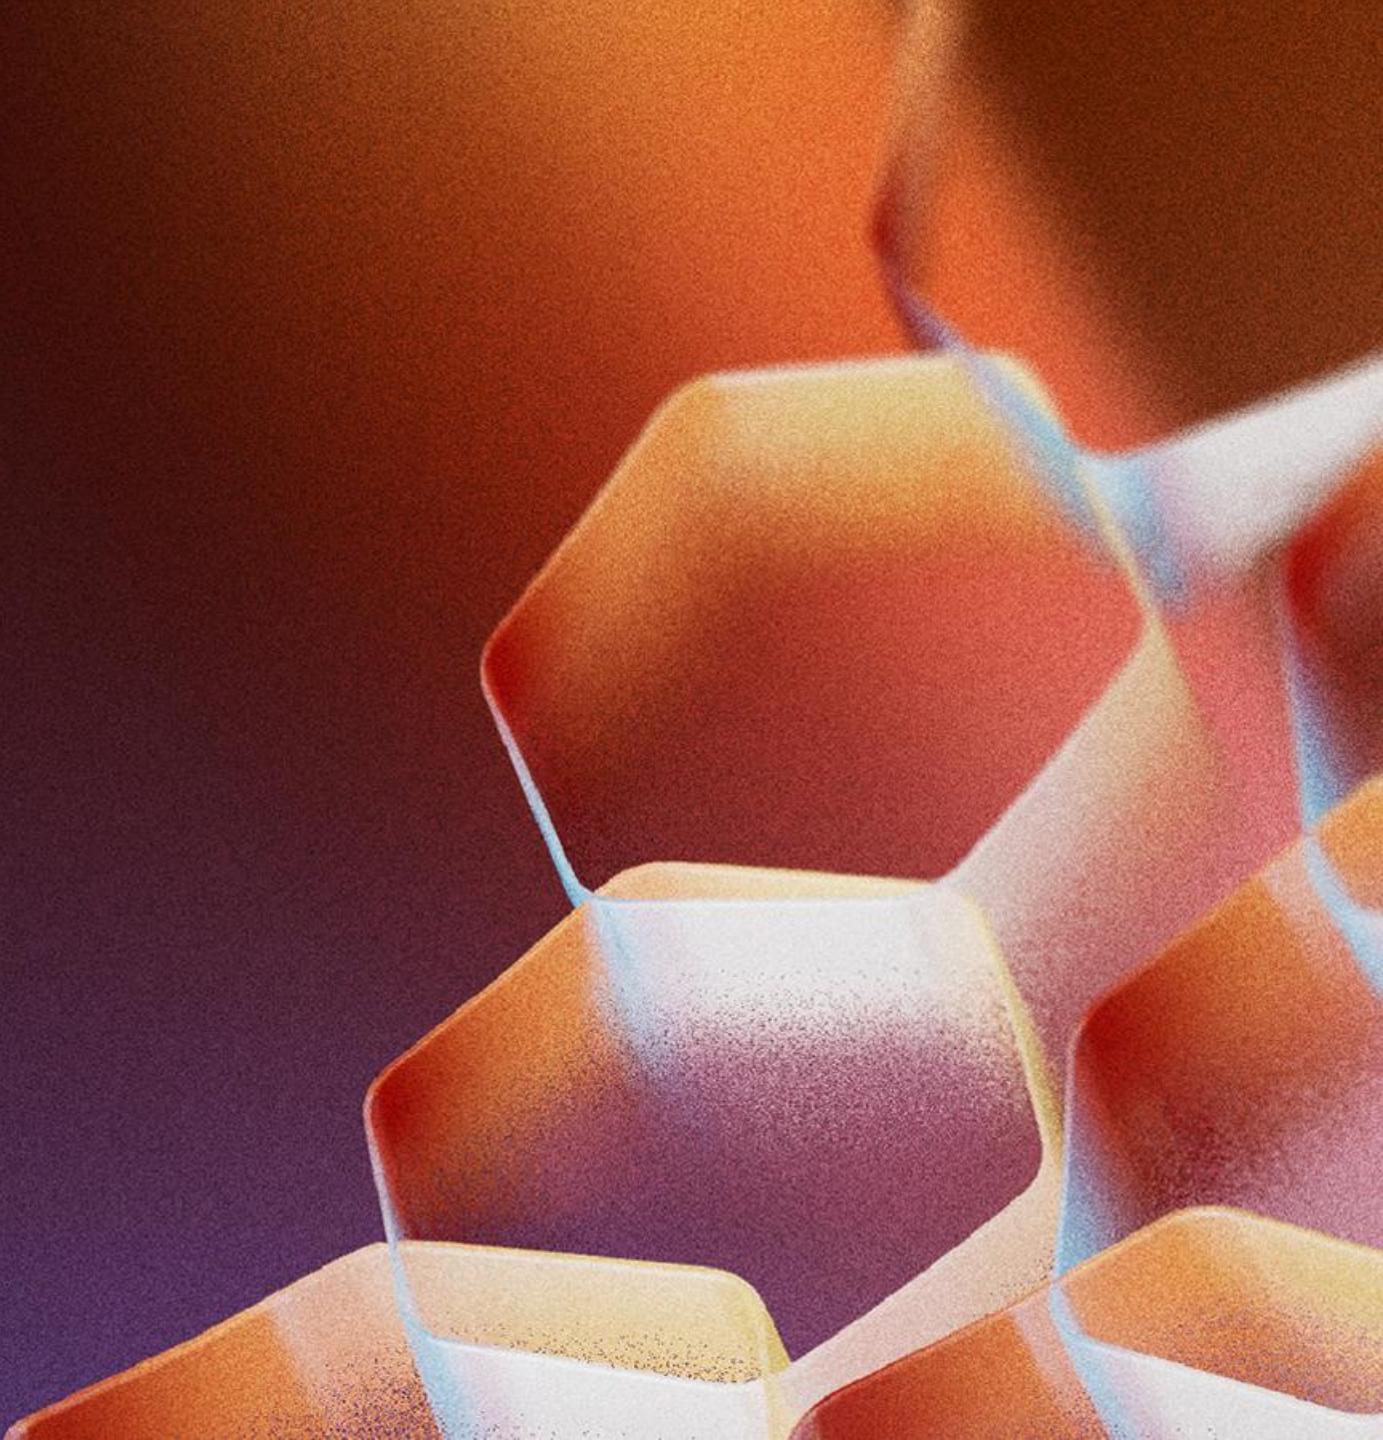
[BRK141: Knowledge Retrieval: RAG for enterprise agents with Azure AI Search](#) (Monday, May 19 6:00PM – 7:00PM)

[BRK142: Agentic RAG: build a reasoning retrieval engine with Azure AI Search](#) (Wednesday, May 21 12:30PM – 1:30PM)



# Model Customization

Nandini Muralidharan



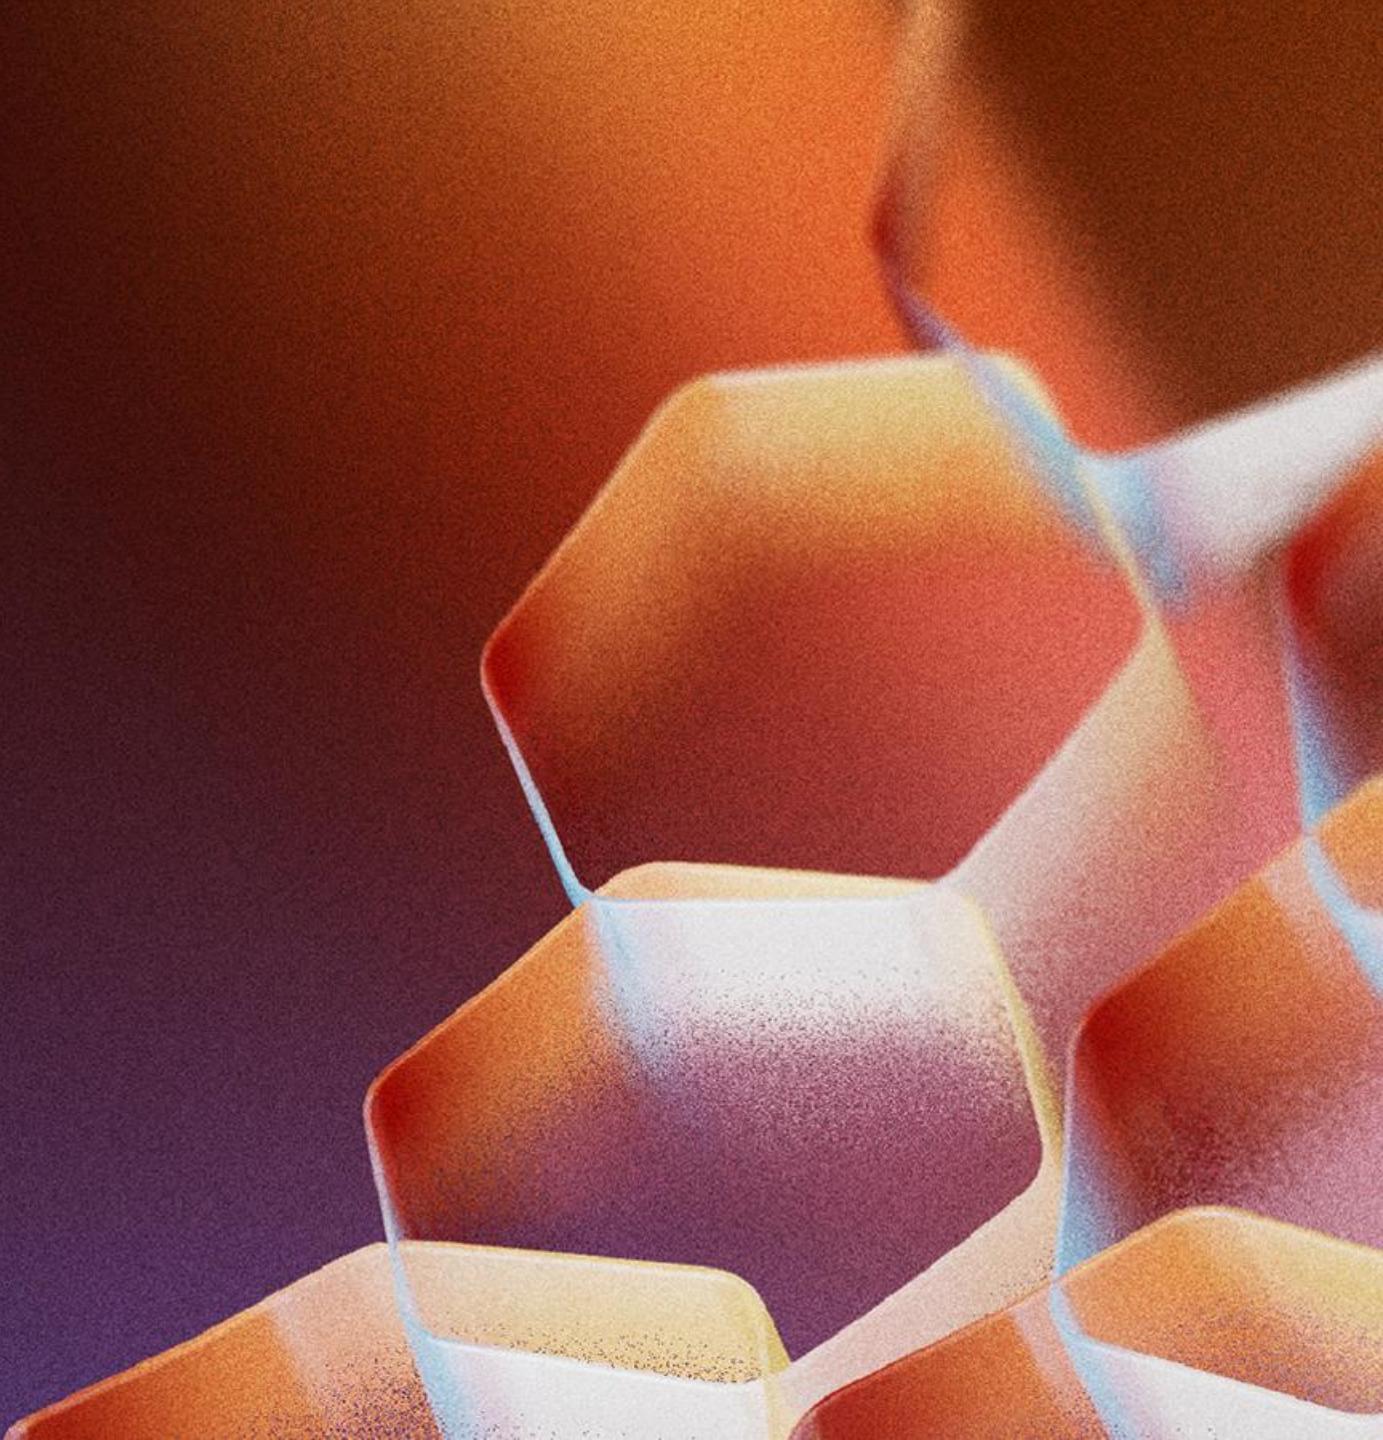
# Model Customization

Tier 2 & 3 Announcements	Content (avail on GH TBA)	Sessions
Models <ul style="list-style-type: none"><li>• SFT – 4.1, 4.1mini, 4.1nano, Llama 4</li><li>• RFT – o4-mini</li></ul>	Models <ul style="list-style-type: none"><li>- SFT : Blog , Distillation demo, Vision demo</li><li>- RFT : Blog, Video, Demo</li></ul>	<a href="#">BRK 150: Fine-tuning and distillation with Azure AI Foundry</a> <a href="#">BRK 155 : Azure AI Foundry: The Agent Factory</a>
Feature enhancements <ul style="list-style-type: none"><li>• Dev Tier</li><li>• Global Training</li><li>• Eval API</li><li>• Cross region deployments</li></ul>	Feature Enhancements <ul style="list-style-type: none"><li>- TC Blog</li><li>- UI</li><li>- TC Blog</li><li>- UI</li></ul>	<a href="#">Fine-Tune End-to-End Distillation Models with Azure AI Foundry</a> <a href="#">Models</a>



# Foundry Local

Jhoanna Machado



# Foundry Local @ Build 2025

## Introducing Foundry Local: Bringing Foundry to Local Devices (In Public Preview)

- KEY010-Keynote: Microsoft Build opening keynote

[Microsoft Build opening keynote](#) (Foundry Local will be mentioned during the keynote)

- BRK146- Breakout: AI on your terms: Build & Ship AI Apps Across Platforms with Windows + Foundry

[Bring AI Foundry to Local: Building cutting-edge on-device AI experiences](#) (Breakout session)

- BRK155-Breakout: Azure AI Foundry: The Agent Factory

[Azure AI Foundry: The Agent Factory](#) (Breakout session where Foundry Local will be mentioned)

- DEM526-Build and ship cross platform apps on-device with AI Foundry and ONNX

[Build and ship cross platform apps on-device with AI Foundry and ONNX](#) (Demo session)

# Foundry Local

Tier 1 & 2 Announcements	KBYG (avail on 05/18)	Content (avail on GH TBA)
<ul style="list-style-type: none"><li>• Foundry Local in Public Preview<ul style="list-style-type: none"><li>• Windows and Mac</li><li>• CPUs, GPUs and NPUs</li><li>• Foundry models: DeepSeek R1, Qwen 2.5 Instruct, Phi-4 Reasoning, Mistral and additional ONNX-format models from Hugging Face</li><li>• Foundry CLI to manage local models, tools and agents</li></ul></li></ul>	<ul style="list-style-type: none"><li>• Read me Guide</li></ul>	<ul style="list-style-type: none"><li>• Running deck for the station</li><li>• Product Demos (ppt with read me guidance and link to demos)</li><li>• Videos with Customer and Partner use cases</li><li>• Preview customers/partners testimonials</li><li>• Link to announcement block [05/19 after 10AM local time]</li></ul>



# Azure AI Document Intelligence

Aditi Maheshwari

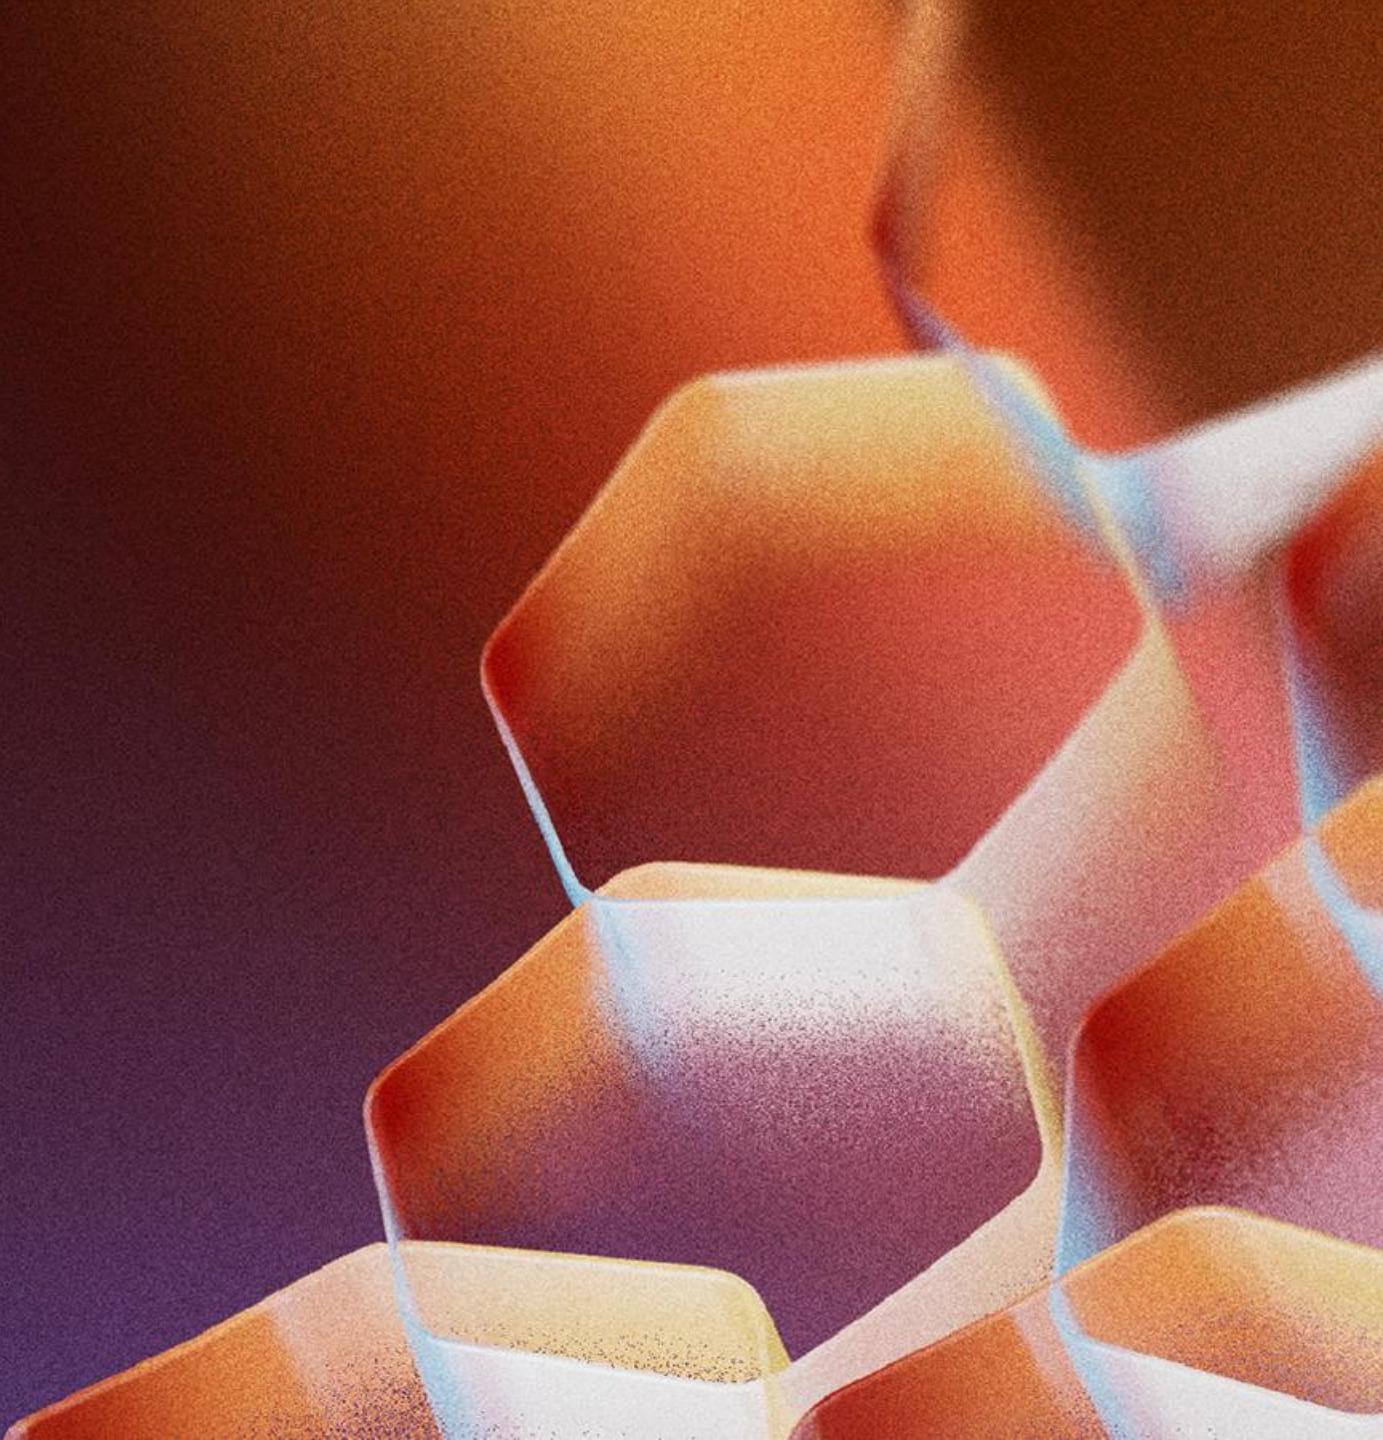
# Azure AI Document Intelligence/ Content Understanding

Tier 1 & 2 Announcements	Content (avail on GH TBA)	Sessions
<ul style="list-style-type: none"><li>• Preview.2 API launch for Content Understanding</li><li>• New Modes (Standard vs Pro – <b>reasoning mode</b>)</li><li>• Classification API</li><li>• Documents – Extended file support(.rtf,.txt,.xml,.json,.msg,.eml,.csv, and .xlsx.) , cross page table, selection mark, confidence score improvements</li><li>• Video – Whole fields support , segmentation support, face identification, face blurring</li><li>• Audio-Multi speaker detection</li><li>• Face- Detection, recognition, and enrollment of faces.</li><li>• Billing- Lower cost with new preview API</li></ul>	<ul style="list-style-type: none"><li>• Announcement slides</li><li>• Demos</li><li>• Videos</li><li>• Blogs</li><li>• FAQ</li></ul>	<a href="#"><u>BRK:169 Multimodal processing for agents with Azure AI Content Understanding</u></a>



# Azure AI Tools

JJ Kim



# Azure AI Tools

Tier 1 & 2 Announcements	Content (avail on GH TBA)	Sessions
<ul style="list-style-type: none"><li>• T2 – PuPr: Offer both Nmt And Llm Translation Through Translator API, 5/19</li><li>• T2 – PuPr: Voice Live API, 5/20</li><li>• T2 – GA: Video Translation Service With New Voice Conversion And Lip Sync, 5/20</li><li>• T2 – PuPr: Embedded Real Time Speech Translation, 5/20</li><li>• T2 – GA: Adaptive Customized Translation, 5/20</li></ul>	<ul style="list-style-type: none"><li>• Announcement slides</li><li>• Demos</li><li>• Videos</li></ul>	<p><a href="#">BRK144 Transform your agentic apps with voice in Foundry</a></p> <p><a href="#">BRK169 Reasoning on multimodal content for efficient agentic AI app building</a></p>



# Azure AI Foundry Content Safety

Ying Chang

# Azure AI Foundry Content Safety

## Sessions

From code to cloud: new tools for building trustworthy GenAI apps

New releases	Release state	Description	Documentation
PII	Public Preview	Configure your LLM output to detect and block over 40 types of personally identifiable information.	in progress (will link to public docs page when ready) (internal link: <a href="#">Personally Identifiable Information (PII) Filter - Azure AI services</a>   Microsoft Learn)
XPIA spotlighting	Public Preview	Take an extra step to defend against prompt injection attacks by "spotlighting" your uploaded documents, which makes it less likely for LLMs to interpret them as instructions.	in progress (will link to public docs page when ready) (internal link: <a href="#">Content Filter Prompt Shields - Azure AI services</a>   Microsoft Learn)
Task adherence monitoring		Ensure your agents are aligned with user intent. The Task Adherence API (coming soon) detects misaligned tool requests that can be blocked or transitioned to Human-In-The-Loop.	No docs yet
Security alerts and recommendations	Public preview	Security alerts and recommendations are now in Foundry for AI developers to monitor and strengthen security posture.	No docs yet
AOAI Request-Level Content Filters	Public preview	Specify your custom AOAI content filtering configuration for every API call at request time.	<a href="#">Use content filters (preview) - Azure OpenAI</a>   Microsoft Learn
Build trustworthy AI with additional security and safety features in Foundry	GA	Enhance the safety and security of your AI deployments with better monitoring, agent trust, guardrails and governance in place.	



# Azure AI Foundry Observability

Sebastian Kohlmeier

## Sessions

[AI and Agent Observability in Azure AI Foundry and Azure Monitor](#)

[Continuously improve your Agent in production](#)

New releases	Release state	Description	Documentation
Enhanced AI Governance	Public Preview	<p>AI governance integrations in Azure AI Foundry, enabling seamless interoperability with Microsoft Purview Compliance Manager, Credo AI and Saidot. These integrations let governance teams define evaluation plans aligned with regulatory frameworks like the EU AI Act and map them directly to the Azure AI Evaluation SDK for execution—eliminating guesswork. Developers can run standardized evaluations using prebuilt notebooks, with results flowing back into governance platforms to support audit readiness and collaboration. Foundry makes AI governance a native, streamlined, and scalable part of the development lifecycle.</p>	
Leaderboard and Benchmarks	GA	<p>Backed by industry-standard benchmarks, these leaderboards help you confidently identify the most suitable foundation models for your AI solutions. Azure AI Foundry Leaderboards support multiple comparison views:</p> <ul style="list-style-type: none"><li>• <b>Quality, cost, and performance leaderboards</b> for quick insights across single metrics<ul style="list-style-type: none"><li>◦ <i>Quality benchmarks</i> assess accuracy on standard datasets across reasoning, knowledge, question answering, math, and coding tasks</li><li>◦ <i>Performance benchmarks</i> track latency, throughput, and time to first token</li><li>◦ <i>Cost benchmarks</i> estimate deployment and usage expenses</li></ul></li><li>• <b>Trade-off charts</b> to visualize how model capabilities (quality) perform under constraints like safety, cost, or throughput</li><li>• <b>Scenario-based leaderboards</b> to surface top-performing models for specific use cases</li><li>• <b>[Pre-announce] Safety benchmarks:</b> Our latest safety benchmarks measure model vulnerability to domain-specific attacks in biosecurity, cybersecurity, and chemical security, using the Attack Success Rate metric. The Weapons of Mass Destruction Proxy (WMDP) benchmark evaluates hazardous knowledge exposure in models—lower Attack Success Rates signal safer model behavior.</li></ul>	<a href="#">Model leaderboards in the Azure AI Foundry portal</a>
Evaluate and Debug with Traces in the Agents Playground	Public Preview	<p>Agents Playground now includes built-in evaluation and tracing capabilities, enabling you to gain deeper insights into your Agent's performance and behavior. When testing an Agent in the playground you can now use evaluation results to assess how your Agent is performing:</p> <ul style="list-style-type: none"><li>• Out of the box evaluation: Quality evaluators are enabled by default and safety evaluators can be enabled as needed.</li><li>• Tracing for debugging: Evaluation results are automatically attached to traces to enable detailed debugging and you can click into threads to see the full agentic execution flow including tool calls, input/output and evaluation metrics.</li></ul>	<a href="#">How to enable tracing in Azure AI Agents - Azure AI Foundry   Microsoft Learn</a>
Evaluate and Debug with Agent-Specific Evaluators	Public Preview	<p>We're excited to announce we have extended our comprehensive suite of quality and safety metrics to evaluate directly Azure AI agent thread messages via a converter. These metrics provide deeper insights into agent performance in both quality and safety aspects, including:</p> <ul style="list-style-type: none"><li>• Intent Resolution: Measures whether an agent identifies and understands a user's request.</li><li>• Relevance: Measures whether an agent's final response is relevant to the user request.</li><li>• Task Adherence: Measures whether an agent's response adheres to assigned tasks according to task instructions.</li><li>• Indirect jailbreak: Measures whether an agentic workflow contains indirect jailbreak attack.</li><li>• Code vulnerability: Measures whether an agentic workflow contains malicious code with security vulnerabilities, such as code injection, tar-slip, SQL injections.</li></ul>	<a href="#">Agent Evaluation with Azure AI Evaluation SDK - Azure AI Foundry   Microsoft Learn</a>

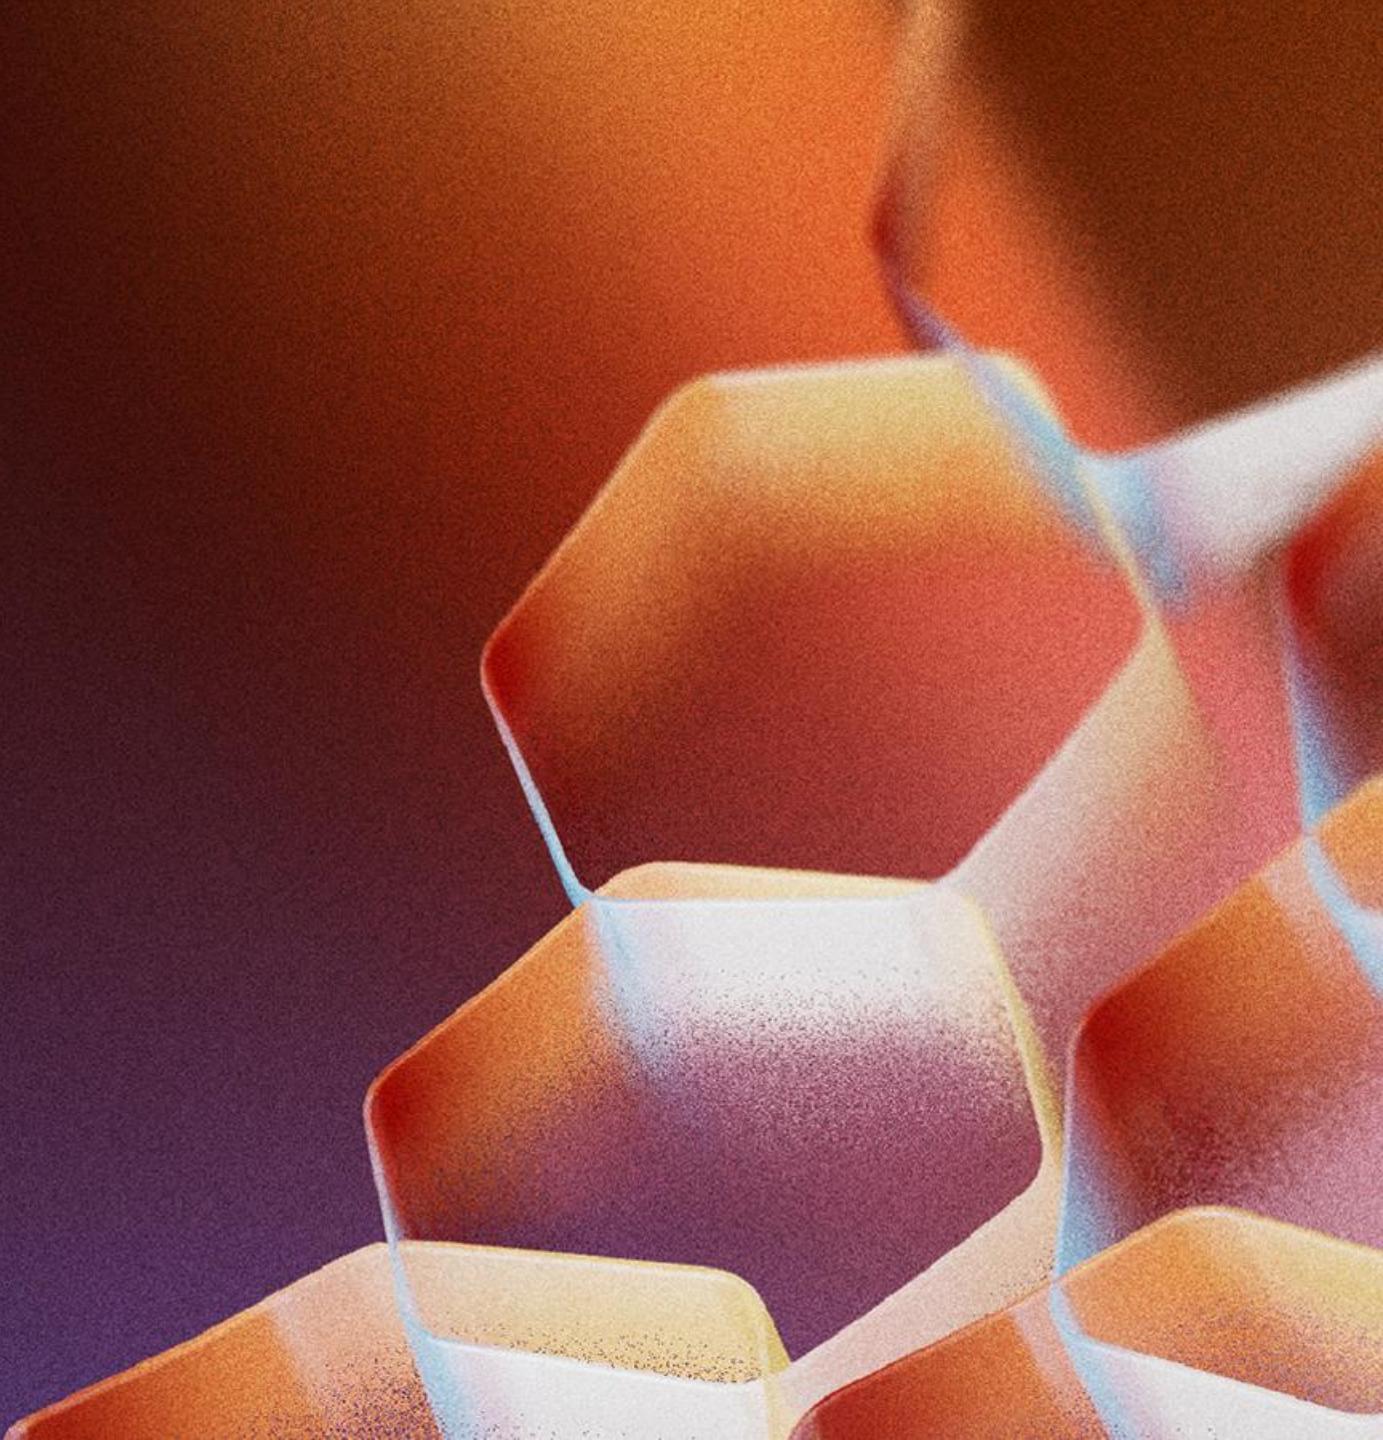
# Azure AI Foundry Observability

New releases	Release state	Description	Documentation
Enhanced Evaluation with Azure OpenAI Graders	Public Preview	<p>We're excited to introduce Azure OpenAI Graders in the Azure AI Foundry SDK, providing comprehensive assessment capabilities for AI models and their outputs. These versatile evaluation tools enable developers to measure performance across multiple dimensions with precision.</p> <p>The suite includes four specialized graders:</p> <ul style="list-style-type: none"><li>• Label Grader: Classifies outputs based on custom-defined labels with detailed explanations</li><li>• String Checker: Compares text to reference values using exact or partial matching</li><li>• Text Similarity: Measures how closely text matches references using metrics like BLEU or ROUGE</li><li>• General Grader: Enables integration of custom graders for specialized evaluation needs</li></ul> <p>Each grader returns detailed results including scores, pass/fail determinations, and explanations. They can be run locally or remotely, offering flexibility in your evaluation workflow.</p>	<a href="#">Azure OpenAI Graders</a>
Azure AI Red Teaming Agent to Scan for Vulnerabilities	Public Preview	<p>Our <a href="#">Azure AI Foundry AI Red Teaming Agent</a> systematically evaluates how your AI models and applications respond to adversarial probing, helping you uncover potential safety vulnerabilities before deployment.</p> <p>The AI Red Teaming Agent integrates Microsoft's open-source <a href="#">PyRIT</a> (Python Risk Identification Tool) into Azure AI Foundry, enabling developers to:</p> <ul style="list-style-type: none"><li>• <b>Automatically scan</b> for content safety risks using simulated adversarial probing</li><li>• <b>Evaluate vulnerability</b> through metrics like Attack Success Rate (ASR)</li><li>• <b>Generate comprehensive reports</b> to assess deployment readiness</li></ul>	<a href="#">Azure AI Foundry AI Red Teaming Agent</a>
CI/CD Workflow Integration	Public Preview	<p>Azure AI Foundry now <b>seamlessly integrates with CI/CD workflows, enabling rapid iteration, evaluation, and scaling</b>. This integration supports continuous development and deployment, ensuring that AI applications and Agents meet quality, performance and safety requirements. By leveraging our CI/CD integrations using the <a href="#">Azure AI Agent Evaluation GitHub Action</a> and Azure Dev Ops Extension, you will be able to run evaluations to assess a single Agent or compare multiple Agent versions every time you iterate on your Agent configuration including the deployed model, model parameters, system prompt and tool settings. These workflows enable you to:</p> <ul style="list-style-type: none"><li>• <b>Automate the pre-production assessment</b> of your Azure AI Agents in your CI/CD workflows using event-based (e.g., on commit), scheduled or manual triggers</li><li>• <b>Evaluate</b> a single Agent or compare multiple Agent versions or Agents.</li><li>• <b>Use built-in quality, performance and safety evaluators</b> from the Azure AI Evaluation SDK.</li><li>• <b>In compare mode</b>, evaluation results include confidence intervals and test for statistical significance so that you can make data-driven decisions with confidence.</li></ul>	<a href="#">Azure AI Agent Evaluation GitHub Action</a> <a href="#">How to run an evaluation in GitHub Action - Azure AI Foundry   Microsoft Learn</a> <a href="#">How to run an evaluation in Azure DevOps - Azure AI Foundry   Microsoft Learn</a>
Continuous Evaluation & Monitoring	Public Preview	<p>Azure AI Foundry enables continuous evaluation and monitoring of agentic systems in production with a new single pane of glass dashboard where you can monitor key metrics including performance, quality, safety, and resource usage. In addition to resource-specific metrics which are available out of the box, we are excited to announce continuous evaluation as a <b>new monitoring capability that runs quality and safety evaluations at a sampled rate of production usage</b> (e.g., 10 evaluations per hour).</p> <p>Once your agent is running in production we provide the following:</p> <ul style="list-style-type: none"><li>• Continuous evaluation and monitoring metrics are made available in the Azure AI Foundry Monitoring dashboard</li><li>• Data is published to Application Insights where you can set up alerts to detect drift or regressions</li><li>• We link from the monitoring dashboard to Azure Monitor so that you can leverage the extended capabilities provided by Azure Monitor to monitor your entire AI application stack.</li><li>• Evaluation results are attached to traces for debugging.</li></ul>	<a href="#">Monitor your Generative AI Applications - Azure AI Foundry   Microsoft Learn</a> <a href="#">Continuously Evaluate your AI agents - Azure AI Foundry   Microsoft Learn</a> <a href="#">How to trace your application with Azure AI Inference SDK - Azure AI Foundry   Microsoft Learn</a>



# AI Skilling

Lanna Teh



# AI Skilling @ Build | strategy



## Overview

Increase awareness and engagement for skilling offerings by connecting with new and existing audiences with key resources, while showcasing our thought leadership, initiatives, and announcements cohesively throughout event.



## Key messages

- **AI Skilling:** Microsoft Learn is your trusted source to help you get skilled up and ready to power AI transformation with the Microsoft Cloud.
- **AI Skills Fest Challenges:** Compete with fellow learners, don't miss the Microsoft AI Skills Fest Challenge running through May 28 to enjoy building new skills in a gamified environment.
- **Plans on Learn:** Leave the event with a curated plan of content on Microsoft Learn with milestones designed to efficiently lead you to upskill on specific AI learning outcomes aligned to Build sessions.
- **Microsoft Credentials:** Microsoft helps make you essential at your current workplace—and the next one with role-based certifications and scenario-specific applied skills that showcase your proficiency with industry-trusted and verified credentials



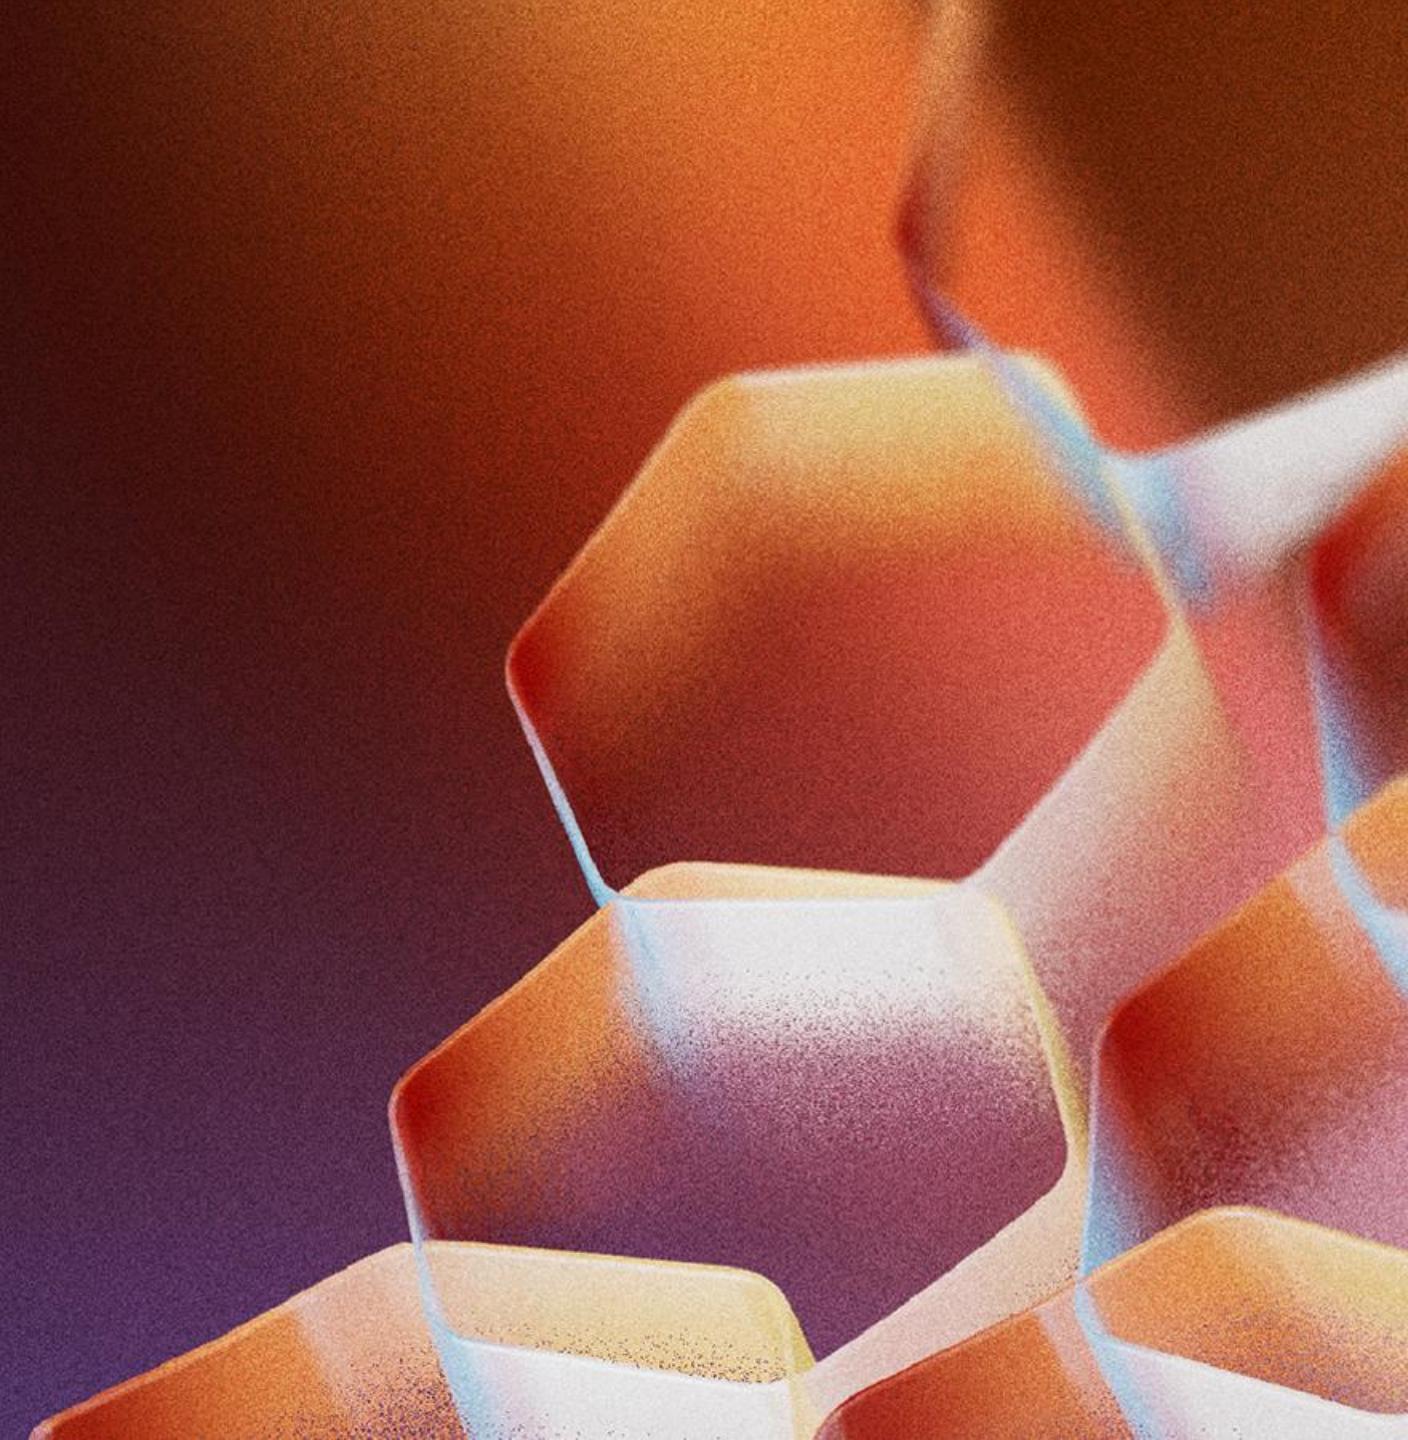
## Key CTAs

- Join an AI Skills Fest Challenge
- Explore Plans on Learn (next step content) aligned to Build sessions
- Validate your skills by earning a Microsoft Credential



# Community

Jonathan Jones



# Hub Overview & Hours



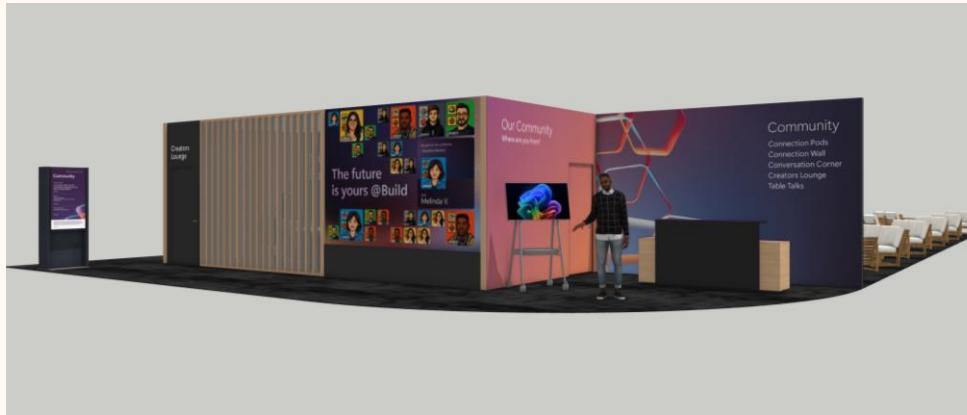
Sunday, May 18	Monday, May 19	Tuesday, May 20	Wednesday, May 21
Early access hours: 9:00 AM – 5:00 PM	Early access hours: 9:00 AM – 10:30 AM	Early access hours: 9:00 AM – 11:00 AM	Early access hours: 8:30 AM – 9:00 AM
	Hub Open: 10:30 AM – 7:30 PM	Hub Open: 11:00 AM – 6:30 PM	Hub Open: 9:00 AM – 6:00 PM

# Community Space Activations

## Creator's Lounge, Connection Wall Check-in Desk



Community experience locations



Left to right: Creators Lounge, Connection Wall, Check-in Desk



Creators Lounge

**Check-in Desk:** Located front-and-center in the Community space, attendees can stop by to learn about Communities or Community content at Microsoft Build, view content schedule for the day.

**Creators Lounge:** This is a branded, reservable space for influencers to livestream and record content for social media.

**Connection Wall:** Digital display to create avatar and connect with attendees based on "developer genetic makeup" into family trees. These commonalities can filter into community sessions to facilitate conversation starters and encourage organic networking.

# Community Space Activations – Content Sessions

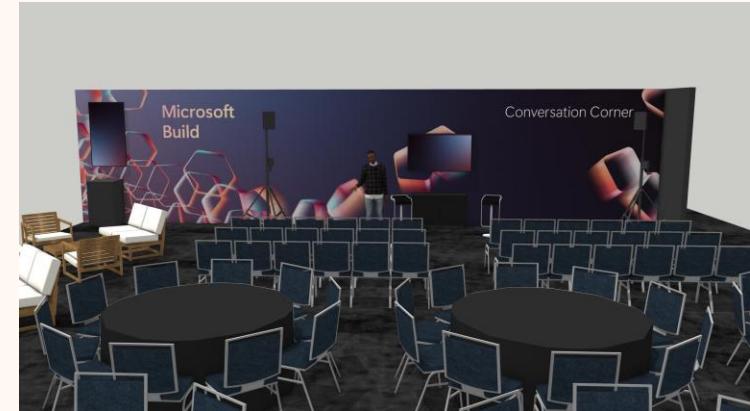
## Conversation Corner & Table Talks



Community experience locations



Conversation Corner



Conversation Corner and Community Table Talks

**Table Talks:** A collection of “table groupings” to support scheduled community table topics. The topics will have a set 3 of tables, each with a designated speaker or table host. 60 minutes in length, roundtables of 10 attendees each.

**Conversation Corner:** A scheduled, dedicated space with “theater-style” seating that can be leveraged for expert content delivery, and one-to-many discussions. 45 minutes in length, seating for 35.

# Community Space Activations

## Activity Walls & Connection Pods



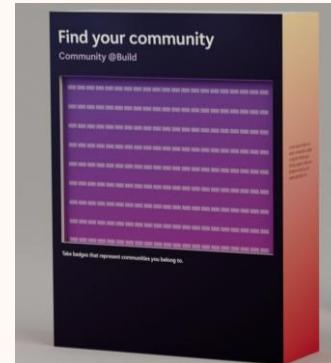
Community experience locations



Light bright wall



Community message wall



Community sticker wall



Connection pod layout

**Activity Walls:** Fun, hands-on engagements to encourage Community sharing and connectivity.

**Connection Pods:** Space designed for "birds of a feather" conversations. These areas facilitate conversations and discussions amongst attendees on pre-determined topics of interest based on community groups.

# AI Community | content programming

## **COMM432: Table Talks: Azure AI Foundry Communities**

Monday, May 19 | 4:00 PM - 5:00 PM Pacific Daylight Time

Speakers: Nick Brady, Carlotta Castelluccio, Korey Stegared-Pace and Laura Wolf

Abstract: Join us to learn more about how to get involved and become an active member of the growing online communities on Discord and GitHub

## **COMM401: Conversations: Ask me anything w/Pablo Castro, CVP of Azure AI Search**

Tuesday, May 20 | 5:45 PM – 6:30 PM Pacific Daylight Time

Speakers: Pablo Castro, Roelant Dieben

Abstract: Join us for an unplugged conversation with Pablo Castro, Corporate Vice President at Microsoft, as we explore the vital role of Azure AI Search in powering generative AI applications. Moderated by AI MVP Roelant Dieben, this interactive session invites you to ask questions and engage in a dynamic discussion. Be sure to attend Pablo's session, *"Knowledge Retrieval: RAG for Enterprise Agents with Azure AI Search,"* on Monday, May 19, to gain foundational knowledge—then bring your follow-up questions to this engaging conversation on Tuesday.

## **COMM433: Table Talks; Azure AI Foundry discussion**

Wednesday, May 21 | 11:00 AM – 12:00 PM Pacific Daylight Time

Speaker: Marco Casalaina

Abstract: Join Marco Casalaina, VP, Core AI at Microsoft, for an open and interactive Table Talk on Azure AI Foundry. This is your opportunity to ask questions, share insights, and explore how Azure AI Foundry is transforming the way organizations build and scale intelligent applications and agents. Whether you're curious about the platform's capabilities or eager to get started, your voice is part of the conversation

## **COMM431: Table Talks Global AI Community**

Wednesday, May 21 | 5:00 PM – 6:00 PM Pacific Daylight Time

Speaker: Henk Boelman, Alexander Wachtel

Abstract: Join this global discussion about community! Discover how to start a community in your area, engage with others, and connect with leaders from around the world.

